

5103 Project 1

Generated by Doxygen 1.7.6.1

Tue Feb 21 2012 15:18:54

Contents

1	Class Index	1
1.1	Class Hierarchy	1
2	Class Index	3
2.1	Class List	3
3	File Index	5
3.1	File List	5
4	Class Documentation	7
4.1	AbstractDevice Class Reference	7
4.1.1	Detailed Description	8
4.1.2	Member Function Documentation	8
4.1.2.1	setTimer	8
4.2	cAbsQueuedDevice Class Reference	8
4.2.1	Member Function Documentation	9
4.2.1.1	scheduleDevice	9
4.2.1.2	setDefaultTime	9
4.2.1.3	timerFinished	9
4.3	cBlockDevice Class Reference	10
4.3.1	Member Function Documentation	11
4.3.1.1	scheduleDevice	11
4.3.1.2	setDefaultTime	12
4.3.1.3	timerFinished	12
4.4	cCharDevice Class Reference	13
4.4.1	Member Function Documentation	14

4.4.1.1	scheduleDevice	14
4.4.1.2	setDefaultTime	14
4.4.1.3	timerFinished	15
4.5	cCPU Class Reference	15
4.5.1	Detailed Description	17
4.5.2	Member Function Documentation	17
4.5.2.1	executePrivSet	17
4.5.2.2	getOpcode	17
4.5.2.3	getParam	17
4.5.2.4	getPSW	17
4.5.2.5	getSetPC	18
4.5.2.6	getSetPSW	18
4.5.2.7	getSetVC	18
4.5.2.8	run	18
4.5.2.9	setMaxPC	18
4.5.2.10	setPSW	18
4.5.2.11	setText	19
4.5.2.12	setUserMode	19
4.5.3	Member Data Documentation	19
4.5.3.1	KMode	19
4.6	cFCFS Class Reference	19
4.6.1	Member Function Documentation	21
4.6.1.1	addLogger	21
4.6.1.2	addProcess	21
4.6.1.3	addProcLogger	22
4.6.1.4	getNextToRun	22
4.6.1.5	initProcScheduleInfo	23
4.6.1.6	numProcesses	23
4.6.1.7	printUnblocked	23
4.6.1.8	removeProcess	24
4.6.1.9	setBlocked	24
4.6.1.10	unblockProcess	24
4.7	cIDManager Class Reference	25
4.7.1	Constructor & Destructor Documentation	26

4.7.1.1	cIDManager	26
4.7.2	Member Function Documentation	26
4.7.2.1	getID	26
4.7.2.2	getLowID	27
4.7.2.3	nextLowID	27
4.7.2.4	reservedIDs	27
4.7.2.5	returnID	27
4.8	cKernel Class Reference	27
4.8.1	Constructor & Destructor Documentation	29
4.8.1.1	cKernel	29
4.8.2	Member Function Documentation	30
4.8.2.1	boot	30
4.8.2.2	cleanupProcess	30
4.8.2.3	cleanupProcess	30
4.8.2.4	initProcess	30
4.8.2.5	sigHandler	30
4.8.2.6	swapProcesses	31
4.8.3	Friends And Related Function Documentation	31
4.8.3.1	deviceHandle	31
4.9	ClockDevice Class Reference	31
4.9.1	Member Function Documentation	32
4.9.1.1	disarm	33
4.9.1.2	getTime	33
4.9.1.3	setTimer	33
4.10	cLottery Class Reference	33
4.10.1	Member Function Documentation	35
4.10.1.1	addProcess	35
4.10.1.2	getNextToRun	36
4.10.1.3	initProcScheduleInfo	36
4.10.1.4	numProcesses	36
4.10.1.5	removeProcess	37
4.10.1.6	setBlocked	37
4.10.1.7	unlockProcess	38
4.11	cMultiLevel Class Reference	38

4.11.1	Member Function Documentation	40
4.11.1.1	addLogger	40
4.11.1.2	addProcess	41
4.11.1.3	addProcLogger	41
4.11.1.4	getNextToRun	41
4.11.1.5	initProcScheduleInfo	42
4.11.1.6	numProcesses	42
4.11.1.7	printUnblocked	43
4.11.1.8	removeProcess	43
4.11.1.9	setBlocked	43
4.11.1.10	unlockProcess	44
4.12	cProcessLogger Class Reference	44
4.12.1	Detailed Description	46
4.13	cRoundRobin Class Reference	46
4.13.1	Member Function Documentation	48
4.13.1.1	getNextToRun	48
4.13.1.2	removeProcess	48
4.13.1.3	setBlocked	48
4.13.1.4	unlockProcess	48
4.14	cScheduler Class Reference	49
4.14.1	Member Function Documentation	50
4.14.1.1	addLogger	50
4.14.1.2	addProcess	50
4.14.1.3	addProcLogger	50
4.14.1.4	getNextToRun	51
4.14.1.5	initProcScheduleInfo	52
4.14.1.6	numProcesses	52
4.14.1.7	printUnblocked	52
4.14.1.8	removeProcess	52
4.14.1.9	setBlocked	53
4.14.1.10	unlockProcess	53
4.15	fcfsInfo Struct Reference	54
4.16	kernelError Struct Reference	54
4.16.1	Detailed Description	55

4.17	lotteryInfo Struct Reference	55
4.18	ProcessInfo Struct Reference	56
4.18.1	Detailed Description	56
4.18.2	Member Data Documentation	56
4.18.2.1	scheduleData	56
4.19	roundRobinInfo Struct Reference	57
4.20	sMultiInfo Struct Reference	57
5	File Documentation	59
5.1	include/cpu.h File Reference	59
5.1.1	Detailed Description	61
5.1.2	Define Documentation	61
5.1.2.1	MAX_PARAM_SIZE	61
5.1.3	Enumeration Type Documentation	61
5.1.3.1	ePSW	61
5.2	include/devices/block_device.h File Reference	61
5.2.1	Detailed Description	62
5.3	include/devices/char_device.h File Reference	62
5.3.1	Detailed Description	64
5.4	include/devices/clock_device.h File Reference	64
5.4.1	Detailed Description	65
5.5	include/devices/queued_device.h File Reference	65
5.5.1	Detailed Description	66
5.6	include/kernel.h File Reference	66
5.6.1	Detailed Description	68
5.6.2	Define Documentation	68
5.6.2.1	DEFAULT_PRIORITY	68
5.6.3	Variable Documentation	68
5.6.3.1	initProcessName	68
5.6.3.2	procLogFile	68
5.6.3.3	traceLogFile	69
5.7	include/process.h File Reference	69
5.7.1	Detailed Description	70
5.7.2	Enumeration Type Documentation	70

5.7.2.1	eProcState	70
5.8	include/scheduler/fcfs.h File Reference	71
5.8.1	Detailed Description	72
5.9	include/scheduler/lottery.h File Reference	72
5.9.1	Detailed Description	74
5.10	include/scheduler/round_robin.h File Reference	74
5.10.1	Detailed Description	75
5.11	include/utility/id.h File Reference	75
5.11.1	Detailed Description	76
5.12	include/utility/logger.h File Reference	76
5.12.1	Detailed Description	78
5.13	include/utility/process_logger.h File Reference	78
5.13.1	Detailed Description	79
5.13.2	Variable Documentation	79
5.13.2.1	outputFormat	79
5.13.2.2	procNameReq	79
5.13.2.3	requestError	79

Chapter 1

Class Index

1.1 Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:

AbstractDevice	7
ClockDevice	31
cAbsQueuedDevice	8
cBlockDevice	10
cCharDevice	13
cCPU	15
cIDManager	25
cKernel	27
cProcessLogger	44
cScheduler	49
cFCFS	19
cLottery	33
cMultiLevel	38
cRoundRobin	46
fcfsInfo	54
kernelError	54
lotteryInfo	55
ProcessInfo	56
roundRobinInfo	57
sMultiInfo	57

Chapter 2

Class Index

2.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

AbstractDevice	An abstract device class	7
cAbsQueuedDevice	Abstract class for a device which queues requests	8
cBlockDevice	Queued block device	10
cCharDevice	Queued char device	13
cCPU	A class for emulating a simple cpu	15
cFCFS	First-Come-First-Serve Scheduler	19
cIDManager	A class for managing unique IDs	25
cKernel	Core managing class for this simulated OS	27
ClockDevice	Device for generating repeated clock interrupts	31
cLottery	Lottery Scheduler	33
cMultiLevel	Multi-Level Queue Scheduler	38
cProcessLogger	Class specifically for logging process state information	44
cRoundRobin	Round Robin Scheduler	46
cScheduler	Abstract Interface for Schedulers	49

fcfsInfo	Struct containing process info specific for FCFS scheduling	54
kernelError	Struct containing kernel crash information	54
lotteryInfo	Struct containing process info specific for Lottery scheduling	55
ProcessInfo	Structure for containing process state and data	56
roundRobinInfo	Struct containing process info specific for Round-Robin scheduling	57
sMultiInfo	Struct containing process info specific for Multi-Level scheduling	57

Chapter 3

File Index

3.1 File List

Here is a list of all documented files with brief descriptions:

include/ cpu.h	59
include/ init.h	??
include/ kernel.h	66
include/ process.h	69
include/ top.h	??
include/devices/ abstract_device.h	??
include/devices/ block_device.h	61
include/devices/ char_device.h	62
include/devices/ clock_device.h	64
include/devices/ queued_device.h	65
include/scheduler/ allSchedulers.h	??
include/scheduler/ fcfs.h	71
include/scheduler/ lottery.h	72
include/scheduler/ multi_level.h	??
include/scheduler/ round_robin.h	74
include/scheduler/ scheduler.h	??
include/utility/ id.h	75
include/utility/ logger.h	76
include/utility/ process_logger.h	78

Chapter 4


Class Documentation

4.1 AbstractDevice Class Reference

An abstract device class.

```
#include <abstract_device.h>
```

Inheritance diagram for AbstractDevice:



Public Member Functions

- virtual void **setTimer** (int time)=0
Set the length of the timer for this device.
- virtual void **disarm** ()=0
*Disarm the timer set from a previous call to **setTimer**.*

4.1.1 Detailed Description

This was initially going to be the abstract class for all devices such as the clock, character and block devices but it turned out this was not the desired interface for everything. This is why the [cAbsQueuedDevice](#) class was made. This class remains here because it is used by the clock device however it is no longer necessary.

4.1.2 Member Function Documentation

4.1.2.1 `void AbstractDevice::setTimer (int time)` `[pure virtual]`

Parameters

<i>int</i>	time It is up to the implementation what the scale for this time is. It will likely be milliseconds or more.
------------	--

Implemented in [ClockDevice](#).

The documentation for this class was generated from the following file:


- `include/devices/abstract_device.h`

4.2 cAbsQueuedDevice Class Reference

Abstract class for a device which queues requests.

```
#include <queued_device.h>
```

Inheritance diagram for `cAbsQueuedDevice`:



Public Member Functions

- virtual void [setDefaultTime](#) (int usec)=0

Set the default timer length for this device.

- virtual void [scheduleDevice](#) ([ProcessInfo](#) *)=0

Schedule a process for a device interrupt.

- virtual [ProcessInfo](#) * [timerFinished](#) ()=0

Called when a device timer has gone off.

- virtual int [queueLength](#) ()=0

Returns the number of devices waiting on the device.

4.2.1 Member Function Documentation

4.2.1.1 void [cAbsQueuedDevice::scheduleDevice](#) ([ProcessInfo](#) *) [pure virtual]

If there are no other waiting processes then schedule the interrupt. If there is a pending interrupt then queue the process.

Parameters

<i>Process-Info*</i>	Process that wants to block on the device. This must be saved and returned to the kernel when its interrupt has been received.
----------------------	--

Warning

This must be synchronized with [timerFinished](#) because signal handlers will cause these methods to be called asynchronously.

Implemented in [cCharDevice](#), and [cBlockDevice](#).

4.2.1.2 void [cAbsQueuedDevice::setDefaultTime](#) (int *usec*) [pure virtual]

Subsequent calls to schedule the device for a process should use this default time. - This is convenience function. A class defined default can be defined statically or in the constructor.

Parameters

<i>int</i>	usec default microsecond length for device timer.
------------	---

Implemented in [cCharDevice](#), and [cBlockDevice](#).

4.2.1.3 [ProcessInfo](#) * [cAbsQueuedDevice::timerFinished](#) () [pure virtual]

The signal handler in the kernel will call this method when a device's timer has completed. The method should then return the waiting process and then schedule an interrupt for the next device queued, if any.

Returns

ProcessInfo* Process that was waiting for the device I/O to complete. It will now be unblocked by the kernel/scheduler.

Warning

Must be synchronized with [scheduleDevice](#). Failing to do so could leave a process blocked indefinitely.

Implemented in [cCharDevice](#), and [cBlockDevice](#).

The documentation for this class was generated from the following file:


- [include/devices/queued_device.h](#)

4.3 cBlockDevice Class Reference


Queued block device.

```
#include <block_device.h>
```

Inheritance diagram for cBlockDevice:



Collaboration diagram for cBlockDevice:



Public Member Functions

- **cBlockDevice** (int usec)
- void **setDefaultTime** (int usec)
Set the default timer length for this device.
- void **scheduleDevice** (**ProcessInfo** *)
Schedule a process for a device interrupt.
- **ProcessInfo** * **timerFinished** ()
Called when a device timer has gone off.
- int **queueLength** ()
Returns the number of devices waiting on the device.

Private Attributes

- queue< **ProcessInfo** * > **waitQueue**
- pthread_mutex_t **deviceLock**

4.3.1 Member Function Documentation

4.3.1.1 void cBlockDevice::scheduleDevice (**ProcessInfo** *) [virtual]

If there are no other waiting processes then schedule the interrupt. If there is a pending interrupt then queue the process.

Parameters

<i>Process-Info*</i>	Process that wants to block on the device. This must be saved and returned to the kernel when its interrupt has been received.
----------------------	--

Warning

This must be synchronized with [timerFinished](#) because signal handlers will cause these methods to be called asynchronously.

Implements [cAbsQueuedDevice](#).

Referenced by `cKernel::boot()`.

4.3.1.2 void **cBlockDevice::setDefaultTime** (int *usec*) [virtual]

Subsequent calls to schedule the device for a process should use this default time. - This is convenience function. A class defined default can be defined statically or in the constructor.

Parameters

<i>int</i>	usec default microsecond length for device timer.
------------	---

Implements [cAbsQueuedDevice](#).

4.3.1.3 **ProcessInfo* cBlockDevice::timerFinished** () [virtual]

The signal handler in the kernel will call this method when a device's timer has completed. The method should then return the waiting process and then schedule an interrupt for the next device queued, if any.

Returns

ProcessInfo* Process that was waiting for the device I/O to complete. It will now be unblocked by the kernel/scheduler.

Warning

Must be synchronized with [scheduleDevice](#). Failing to do so could leave a process blocked indefinitely.

Implements [cAbsQueuedDevice](#).

The documentation for this class was generated from the following files:


- include/devices/[block_device.h](#)
- src/devices/block_device.cpp

4.4 cCharDevice Class Reference

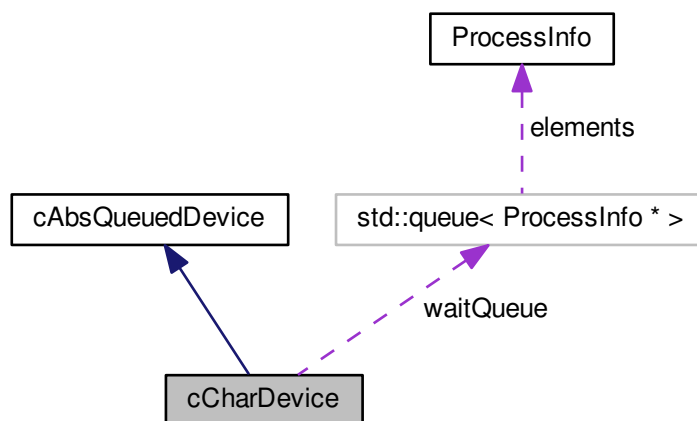
Queued char device.

```
#include <char_device.h>
```

Inheritance diagram for cCharDevice:



Collaboration diagram for cCharDevice:



Public Member Functions

- **cCharDevice** (int usec)

- void [setDefaultTime](#) (int usec)
Set the default timer length for this device.
- void [scheduleDevice](#) ([ProcessInfo](#) *)
Schedule a process for a device interrupt.
- [ProcessInfo](#) * [timerFinished](#) ()
Called when a device timer has gone off.
- int [queueLength](#) ()
Returns the number of devices waiting on the device.

Private Attributes

- queue< [ProcessInfo](#) * > **waitQueue**
- pthread_mutex_t **deviceLock**

4.4.1 Member Function Documentation

4.4.1.1 void [cCharDevice::scheduleDevice](#) ([ProcessInfo](#) *) [virtual]

If there are no other waiting processes then schedule the interrupt. If there is a pending interrupt then queue the process.

Parameters

<i>Process-Info*</i>	Process that wants to block on the device. This must be saved and returned to the kernel when its interrupt has been received.
----------------------	--

Warning

This must be synchronized with [timerFinished](#) because signal handlers will cause these methods to be called asynchronously.

Implements [cAbsQueuedDevice](#).

Referenced by [cKernel::boot\(\)](#).

4.4.1.2 void [cCharDevice::setDefaultTime](#) (int *usec*) [virtual]

Subsequent calls to schedule the device for a process should use this default time. - This is convenience function. A class defined default can be defined statically or in the constructor.

Parameters

<i>int</i>	usec default microsecond length for device timer.
------------	---

Implements [cAbsQueuedDevice](#).

4.4.1.3 ProcessInfo* cCharDevice::timerFinished () [virtual]

The signal handler in the kernel will call this method when a device's timer has completed. The method should then return the waiting process and then schedule an interrupt for the next device queued, if any.

Returns

ProcessInfo* Process that was waiting for the device I/O to complete. It will now be unblocked by the kernel/scheduler.

Warning

Must be synchronized with [scheduleDevice](#). Failing to do so could leave a process blocked indefinitely.

Implements [cAbsQueuedDevice](#).

The documentation for this class was generated from the following files:

- include/devices/[char_device.h](#)
- src/devices/[char_device.cpp](#)

4.5 cCPU Class Reference

A class for emulating a simple cpu.

```
#include <cpu.h>
```

Public Member Functions

- void **initTraceLog** ()
- void **initClockPulse** (pthread_mutex_t *_pulseLock, pthread_cond_t *_pulseCond)
- void **setText** (char *text)
Set the program text.
- void **setMaxPC** (unsigned int newMax)
Set the max PC value.
- void **setUserMode** ()
Set the cpu back into user mode.
- unsigned int **getSetPC** (unsigned int newPC)
Get/Set the program counter.
- int **getSetVC** (int newVC)
Get/Set the VC.
- uint16_t **getSetPSW** (uint16_t newPSW)
Get/Set the PSW.

- uint16_t [getPSW](#) ()
Get the Program Status Word.
- void [setPSW](#) (uint16_t newPSW)
Set a new value for the PSW.
- char * [getParam](#) (int num)
Get execution parameters from the cpu.
- char [getOpcode](#) ()
Get the current Opcode.
- void [run](#) ()
Start execution.
- void [executePrivSet](#) (int num, int &clockTick)
Execute set number of privileged instructions.

Public Attributes

- pidType **pid**

Private Member Functions

- int [tokenizeLine](#) ()

Private Attributes

- int **clockTick**
- bool [KMode](#)
Kernel Mode bit.
- unsigned int **PC**
- unsigned int **maxPC**
- int **VC**
- uint16_t [PSW](#)
Program Status Word.
- char * [execText](#)
Text data for currently executing process.
- char [tokenBuffer](#) [2][MAX_PARAM_SIZE]
Holds the tokenized execution parameters.
- char [Opcode](#)
Holds the current Opcode.
- FILE * **traceStream**
- pthread_mutex_t * [pulseLock](#)
Used when executing privileged set to synchronize with clock.
- pthread_cond_t * [pulseCond](#)
Used in conjunction with [pulseLock](#).

4.5.1 Detailed Description

This class emulates the internals of a very simple cpu with two main registers, PC and VC. In addition, it has other state for handling system calls and program exceptions.

4.5.2 Member Function Documentation

4.5.2.1 void cCPU::executePrivSet (int *num*, int & *clockTick*)

When the kernel receives a syscall for device I/O it will call this function to "simulate" executing kernel mode ops.

Parameters

<i>int</i>	num Number of privileged instructions to execute.
<i>int&</i>	clockTick Reference to the clockTick counter in the kernel.

Referenced by cKernel::boot().

4.5.2.2 char cCPU::getOpcode ()

Get the current Opcode in the cpu. This is used by the kernel to determine which system call is being made. Used in conjunction with [cCPU::getParam](#) the kernel can process system calls.

Referenced by cKernel::boot().

4.5.2.3 char * cCPU::getParam (int *num*)

Fetch the given execution parameter from the cpu's internal buffer. When an instruction is encountered that has parameters associated with it, the cpu tokenizes them and places it in an internal buffer. This function is mainly used by the kernel in handling system calls.

Parameters

<i>num</i>	Must be less than MAX_PARAMS (currently 2)
------------	--

Returns

Returns a char* which points to a string of at most MAX_PARAM_SIZE - 1 bytes.

Referenced by cKernel::boot().

4.5.2.4 uint16_t cCPU::getPSW ()

Returns the program status word which is an unsigned 16-bit integer type with flags from [ePSW](#) set. These are used by the kernel to make action decisions.

Referenced by `cKernel::boot()`.

4.5.2.5 `unsigned int cCPU::getSetPC (unsigned int newPC)`

Get the current value for the program counter and then set its value to the given parameter. This is useful for swapping out process values.

Referenced by `cKernel::swapProcesses()`.

4.5.2.6 `uint16_t cCPU::getSetPSW (uint16_t newPSW)`

Get the current value for the PSW and set its value to the given parameter.

Referenced by `cKernel::swapProcesses()`.

4.5.2.7 `int cCPU::getSetVC (int newVC)`

Get the current value for VC and set its value to the given parameter. This is useful for swapping out process values.

Referenced by `cKernel::boot()`, and `cKernel::swapProcesses()`.

4.5.2.8 `void cCPU::run ()`

Once all appropriate process data is entered by the kernel this function is called to start execution. Any time control needs to be returned to the kernel this function will return with the appropriate PSW flags set for the kernel to act on.

Referenced by `cKernel::boot()`.

4.5.2.9 `void cCPU::setMaxPC (unsigned int newMax)`

This sets the maxPC 'register' in the cpu. This is used when parsing commands to ensure that the cpu does not fall off the end of the process' text segment.

Referenced by `cKernel::swapProcesses()`.

4.5.2.10 `void cCPU::setPSW (uint16_t newPSW)`

Used by the kernel to reset the PSW after a system call. Any process execution which returns to the kernel but does not terminate the process should reset the PSW so subsequent exceptions/terminations are not lost by stray PSW values.

Referenced by `cKernel::boot()`.

4.5.2.11 void cCPU::setText (char * *text*)

Point the cpu to the text data for the running process. This text is indexed using the program counter (PC).

Parameters

<i>text</i>	Program text pointer. assert(text != NULL)
-------------	---

Referenced by cKernel::swapProcesses().

4.5.2.12 void cCPU::setUserMode ()

This is used by the kernel after the kernel has finished servicing a process' kernel mode request (syscall).

Referenced by cKernel::swapProcesses().

4.5.3 Member Data Documentation**4.5.3.1 bool cCPU::KMode [private]**

Set upon system calls.

Referenced by executePrivSet(), run(), and setUserMode().

The documentation for this class was generated from the following files:


- include/cpu.h
- src/cpu.cpp

4.6 cFCFS Class Reference

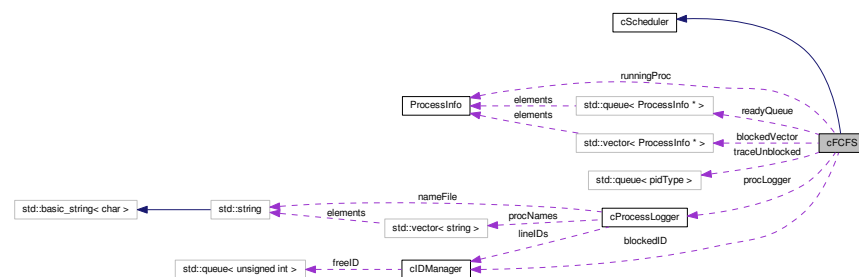
First-Come-First-Serve Scheduler.

```
#include <fcfs.h>
```

Inheritance diagram for cFCFS:



Collaboration diagram for cFCFS:



Public Member Functions

- void **initProcScheduleInfo** ([ProcessInfo](#) *)
Initialize scheduler specific information within the [ProcessInfo](#) struct.
- void **addProcess** ([ProcessInfo](#) *)
Transfer control of process state and scheduling.
- void **setBlocked** ([ProcessInfo](#) *)
Set a process into a *blocked* state.
- void **unlockProcess** ([ProcessInfo](#) *)
Unblock a process and make it ready.
- void **removeProcess** ([ProcessInfo](#) *)
Remove a process from the control of the scheduler.
- [ProcessInfo](#) * **getNextToRun** ()
Query the scheduler for next process to run.
- pidType **numProcesses** ()

How many processes are in the scheduler.

- void [addLogger](#) (FILE *_logStream)

Add an instance of the trace logger to the scheduler.

- void [addProcLogger](#) ([cProcessLogger](#) *_procLogger)

Add instance of process logger to the scheduler.

- void [printUnblocked](#) ()

Prints at once all processes that may have become unblocked asynchronously.

Private Attributes

- queue< [ProcessInfo](#) * > **readyQueue**
- vector< [ProcessInfo](#) * > **blockedVector**
- queue< pidType > **traceUnblocked**
- int **totalBlocked**
- [ProcessInfo](#) * **runningProc**
- pthread_mutex_t **blockedLock**
- pthread_cond_t **allBlocked**
- [cIDManager](#) **blockedID**
- FILE * **logStream**
- [cProcessLogger](#) * **procLogger**

4.6.1 Member Function Documentation

4.6.1.1 void cFCFS::addLogger (FILE *) [virtual]

The kernel passes a file stream pointer for the scheduler to print trace info such as blocking/unblocking events.

Parameters

<i>FILE*</i>	File stream for trace log
--------------	---------------------------

See also

[initLog\(const char* filename\)](#)
[getStream\(\)](#)

Implements [cScheduler](#).

4.6.1.2 void cFCFS::addProcess (ProcessInfo *) [virtual]

After this is called, the kernel core no longer keeps track of the given process. Once the process is created and deemed [ready](#) by the kernel it is handed off here. The scheduler is then in charge of state transitions when the kernel gives it appropriate notifications.

Implementation Requirements:

- Store the process in some location. It should be recognized as ready given its location but the datastructures and organization are implementation specific.

Parameters

<i>Process-Info*</i>	Process to add under scheduler's control
----------------------	--

Implements [cScheduler](#).

4.6.1.3 void cFCFS::addProcLogger (cProcessLogger *) [virtual]

The kernel passes this pointer so the scheduler can update process info upon process addition, state change and termination. This updates the information presented to the top process.

Parameters

<i>cProcess-Logger*</i>	Process logger class
-------------------------	----------------------

See also

[cProcessLogger](#)

Implements [cScheduler](#).

4.6.1.4 ProcessInfo * cFCFS::getNextToRun () [virtual]

After this function is called, it should be assumed by any scheduler implementation that the kernel will run the given process (unless otherwise notified). The currently running process should implicitly be considered for running next (again).

If there are processes left but all are blocked. This function should block until it receives a signal that a process is unblocked.

Implementation Requirements:

- Call [printUnblocked\(\)](#) to update the trace log after a clock step.
- Block if there are > 0 processes but none can run.
- A process listed as [running](#) within the scheduler should be treated as ready when this method is called.
- A process returned as the 'next to run' should be marked as being in a [running](#) state before returning.
- If there are no more remaining processes this should return NULL.

Returns

[ProcessInfo](#)* Ready process to run next. May be the same as the currently running one.

Warning

Must be thread safe with block and unblock methods.

Implements [cScheduler](#).

4.6.1.5 void cFCFS::initProcScheduleInfo ([ProcessInfo](#) *) [virtual]

This method is called after the kernel has initialized all process data but before it is marked [ready](#). This gives the scheduler an opportunity to initialize any scheduler specific data and assign it to the [ProcessInfo::scheduleData](#) member.

Implementation Requirements:

- No required actions.

Implements [cScheduler](#).

4.6.1.6 pidType cFCFS::numProcesses () [virtual]

This returns how many processes, both running and blocked, are being handled by the scheduler.

Implementation Requirements:

- Return how many processes, running and blocked, are in the scheduler.

Returns

pidType

Implements [cScheduler](#).

4.6.1.7 void cFCFS::printUnblocked () [virtual]

This method is needed to avoid mixed output in the trace logger.

Implements [cScheduler](#).

Referenced by getNextToRun().

4.6.1.8 void cFCFS::removeProcess (ProcessInfo *) [virtual]

When a process terminates, either through normal means or an exception, the kernel will call this function to release a process from the scheduler's control. The scheduler should clean up any internal state for the process. Deallocation of process resources is left to the kernel.

Implementation Requirements:

- Scheduler should deallocate any resources it assigned to the process within the [ProcessInfo::scheduleData](#) member.
- The Scheduler should remove any pointers to the give process to avoid dereferencing a dead pointer.
- Implementations must not deallocate any memory except that mentioned above. This is handled by the kernel.
- Implementations should mark the process as [terminated](#).

Parameters

<i>Process-Info*</i>	Process to remove from scheduler
----------------------	----------------------------------

Implements [cScheduler](#).

4.6.1.9 void cFCFS::setBlocked (ProcessInfo *) [virtual]

The kernel will call the scheduler with this function when the process has done an operation which causes it to block (l).

Implementation Requirements:

- Process must be marked [blocked](#) and scheduler state should be changed accordingly.
- After this call a process should not be considered for a scheduling decision

Warning

Must be thread safe. Signal handler/s may block during schedule decision.

Implements [cScheduler](#).

4.6.1.10 void cFCFS::unblockProcess (ProcessInfo *) [virtual]

When a process has completed a blocking call the kernel will notify the scheduler that it should be unblocked. This operation should be very fast since it will likely be called from a signal handler.

Implementation Requirements:

- The process must be unblocked and marked [ready](#). It must be available for scheduling with the next call to `::getNextToRun`

Parameters

<i>Process-Info*</i>	Process to unblock
----------------------	--------------------

Warning

Must be thread safe. Signal handler/s may unblock during schedule decision.

Implements [cScheduler](#).

The documentation for this class was generated from the following files:


- `include/scheduler/fcfs.h`
- `src/scheduler/fcfs.cpp`

4.7 cIDManager Class Reference

A class for managing unique IDs.

```
#include <id.h>
```

Collaboration diagram for cIDManager:



Public Member Functions

- [cIDManager](#) (unsigned int startID=0)

Creates a new ID Manager object.

- unsigned int `getID()`

Reserves a unique ID.

- unsigned int `getLowID()`

Get a low ID.

- void `returnID(unsigned int id)`

Returns an ID to the manager.

- unsigned int `nextLowID()`

See what the next low ID would be.

- unsigned int `reservedIDs()`

How many IDs have been given out.

Private Attributes

- queue< unsigned int > `freeID`

Queue of returned IDs.

- unsigned int `baseID`

To prevent the IDs from dropping below when being returned.

- unsigned int `currentID`

Next ID to be given out.

- bool `consumeQueue`

This signals that IDs have reached their max and freeIDs should be used.

4.7.1 Constructor & Destructor Documentation

4.7.1.1 `cIDManager::cIDManager (unsigned int startID = 0)`

Default start ID is 0.

4.7.2 Member Function Documentation

4.7.2.1 unsigned int `cIDManager::getID()`

Unique is in the sense that no one else is currently using it but it may have been used previously. are distributed in increasing order until `UINT_MAX` is reached. After this is reached, IDs are given from the queue of returned IDs. If this queue is empty then an exception is thrown.

Referenced by `getLowID()`, `cKernel::initProcess()`, `cFCFS::setBlocked()`, `cRoundRobin::setBlocked()`, `cLottery::setBlocked()`, and `cMultiLevel::setBlocked()`.

4.7.2.2 unsigned int cIDManager::getLowID ()

When generating process PID's we use the regular getID so that process IDs continue to grow. This choice was mainly to prevent confusion when process 1 terminated and the next one to start had pid = 1. For functions which use vectors and need an ID system, it is more efficient to maintain a smaller window to keep the array small. This method provides this by preferring to return IDs from the freeID queue. Therefore, if there is an ID available in freeID then the total range of IDs will not grow after this function call.

Referenced by cLottery::addProcess(), and cLottery::unblockProcess().

4.7.2.3 unsigned int cIDManager::nextLowID ()

There is not longer any purpose for this function but I left it here for the potential functionality. The intention was to improve performance in the process logger to determine if the next ID would be right after the previous low ID. That way, if we had variable length records we wouldn't have to search from the beginning.

4.7.2.4 unsigned int cIDManager::reservedIDs ()

Returns the number of IDs which have been reserved

4.7.2.5 void cIDManager::returnID (unsigned int id)

If the ID is not equal to the one last given then it is added to a 'free queue'. If it is equal to the last one reserved then the ID counter is simply decremented. If this last case happens, it causes [cIDManager::getID](#) to stop consuming from the queue and return this newly available ID.

Referenced by cKernel::cleanupProcess(), cLottery::removeProcess(), cLottery::setBlocked(), cFCFS::unblockProcess(), cRoundRobin::unblockProcess(), cLottery::unblockProcess(), and cMultiLevel::unblockProcess().

The documentation for this class was generated from the following files:

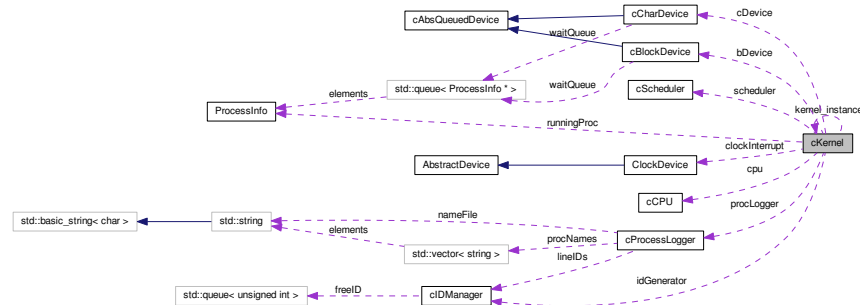
- include/utility/id.h
- src/utility/id.cpp

4.8 cKernel Class Reference

Core managing class for this simulated OS.

```
#include <kernel.h>
```

Collaboration diagram for cKernel:



Public Member Functions

- **cKernel** (**cScheduler** &)
*Default **cKernel** constructor.*
- void **boot** ()
Start the 'OS' Kernel.
- void **initProcess** (const char *filename, pidType parent, int priority=DEFAULT_PRIORITY)
Initialize a Process.
- void **cleanupProcess** (pidType pid)
Cleans up a terminated process.

Private Member Functions

- void **swapProcesses** (**ProcessInfo** *proc, bool switchMode=true)
Swap a process on the cpu.
- void **cleanupProcess** (**ProcessInfo** *)
Cleanup process memory and state.
- void **sigHandler** (int signum, siginfo_t *info)
Handler for all signals.

Static Private Member Functions

- static void **sig_catch** (int signum, siginfo_t *info, void *context)
Static function just for capturing signals.

Private Attributes

- FILE * **traceStream**
- [cProcessLogger](#) **procLogger**
- [cCPU](#) **cpu**
- int **clockTick**
- pthread_mutex_t [intLock](#)
Lock for condition variable.
- pthread_cond_t [intCond](#)
For synchronization with clocktick.
- [cBlockDevice](#) **bDevice**
- [cCharDevice](#) **cDevice**
- [ClockDevice](#) **clockInterrupt**
- pthread_t **deviceThread**
- sem_t [DevSigSem](#)
A device interrupt has been received.
- sem_t [BSigSem](#)
A block device interrupt has been received.
- sem_t [CSigSem](#)
A char device interrupt has been received.
- [ProcessInfo](#) * **runningProc**
Process currently on the cpu.
- [cIDManager](#) **idGenerator**
For generating new process PID's.
- [cScheduler](#) & **scheduler**

Static Private Attributes

- static [cKernel](#) * **kernel_instance**

Friends

- void * [deviceHandle](#) (void *)
Thread function for handline device interrupts.

4.8.1 Constructor & Destructor Documentation

4.8.1.1 cKernel::cKernel (cScheduler & s)

The default constructor initializes all internal datastructures and loads the initial program (default: 'main.trace') but does not run it.

4.8.2 Member Function Documentation

4.8.2.1 void cKernel::boot ()

Starts the main kernel loop. The initial program is loaded and execution follows from there.

Exceptions

kernelError

4.8.2.2 void cKernel::cleanupProcess (ProcessInfo * *proc*) [private]

This is called when a process is being removed from the system and its memory and any remaining state information needs to be cleaned up.

Referenced by boot().

4.8.2.3 void cKernel::cleanupProcess (pidType *pid*)

Cleans up any memory and kernel entries associated with the terminated process. Also removes the process from the scheduler.

4.8.2.4 void cKernel::initProcess (const char * *filename*, pidType *parent*, int *priority* = DEFAULT_PRIORITY)

Initializes a process by loading program file contents, setting default process values and adding it in a ready state to the scheduler.

Referenced by boot(), and cKernel().

4.8.2.5 void cKernel::sigHandler (int *signum*, siginfo_t * *info*) [private]

For clock interrupt signals, the handler signals on a condition variable which the main thread will be waiting on.

For block and character devices, the appropriate semaphores are incremented which will unblock the waiting thread to act on them.

This function is called by [cKernel::sig_catch](#)

See also

[deviceHandle\(void*\)](#)
[cKernel::intCond](#)
[DevSigSem](#)
[CSigSem](#)
[BSigSem](#)

Referenced by `sig_catch()`.

4.8.2.6 `void cKernel::swapProcesses (ProcessInfo * proc, bool switchMode = true)`
[private]

Takes the process in its parameter and swaps it with the one currently running in the cpu.

Referenced by `boot()`.

4.8.3 Friends And Related Function Documentation

4.8.3.1 `void* deviceHandle (void * args)` [friend]

Since limited work can be done within the signal handlers and the main thread may block, a separate thread must be present to handle device interrupts.

See also

[DevSigSem](#)
[CSigSem](#)
[BSigSem](#)

Parameters

<code>void*</code>	This is a pointer to the kernel instance.
--------------------	---

Referenced by `boot()`.

The documentation for this class was generated from the following files:


- `include/kernel.h`
- `src/kernel.cpp`

4.9 ClockDevice Class Reference


Device for generating repeated clock interrupts.

```
#include <clock_device.h>
```

Inheritance diagram for ClockDevice:



Collaboration diagram for ClockDevice:



Public Member Functions

- void `setTimer` (int usec)
Set the timer to go off.
- void `disarm` ()
Disarm the timer.
- int `getTime` ()
Get how much time is remaining.

4.9.1 Member Function Documentation

4.9.1.1 void ClockDevice::disarm () [virtual]

Tries to disarm the timer.

Exceptions

<i>std::string</i>	error message
--------------------	---------------

Implements [AbstractDevice](#).

4.9.1.2 int ClockDevice::getTime ()

Returns the remaining time until a signal is produced.

Returns

int Time left in microseconds

4.9.1.3 void ClockDevice::setTimer (int *usec*) [virtual]

Sets timer to send signal [CLOCKSIG](#) in usec microseconds.

Parameters

<i>usec</i>	Time in microseconds.
-------------	-----------------------

Exceptions

<i>std::string</i>	error message
--------------------	---------------

Implements [AbstractDevice](#).

Referenced by `cKernel::boot()`.

The documentation for this class was generated from the following files:


- `include/devices/clock_device.h`
- `src/devices/clock_device.cpp`

4.10 cLottery Class Reference


Lottery Scheduler.

```
#include <lottery.h>
```

Inheritance diagram for cLottery:



Collaboration diagram for cLottery:



Public Member Functions

- void `initProcScheduleInfo` (`ProcessInfo *`)
Initialize scheduler specific information within the `ProcessInfo` struct.
- void `addProcess` (`ProcessInfo *`)
Transfer control of process state and scheduling.
- void `setBlocked` (`ProcessInfo *`)
Set a process into a `blocked` state.
- void `unlockProcess` (`ProcessInfo *`)
Unblock a process and make it ready.
- void `removeProcess` (`ProcessInfo *`)
Remove a process from the control of the scheduler.
- `ProcessInfo *` `getNextToRun` ()
Query the scheduler for next process to run.
- pidType `numProcesses` ()
How many processes are in the scheduler.

- void [addLogger](#) (FILE * _logStream)
Assign the file pointer to the logStream data field.
- void [addProcLogger](#) ([cProcessLogger](#) * _procLogger)
Assign the cProcessLogger pointer to the procLogger data field.
- void [printUnblocked](#) ()
Removes each PID from the traceUnblocked queue and prints it to the trace logger.

Private Attributes

- vector< [ProcessInfo](#) * > **readyVector**
- vector< [ProcessInfo](#) * > **blockedVector**
- queue< pidType > **traceUnblocked**
- int **totalReady**
- int **totalBlocked**
- int **totalTickets**
- [ProcessInfo](#) * **runningProc**
- pthread_mutex_t **blockedLock**
- pthread_cond_t **allBlocked**
- [cIDManager](#) **blockedID**
- [cIDManager](#) **readyID**
- FILE * **logStream**
- [cProcessLogger](#) * **procLogger**

4.10.1 Member Function Documentation

4.10.1.1 void cLottery::addProcess ([ProcessInfo](#) *) [virtual]

After this is called, the kernel core no longer keeps track of the given process. Once the process is created and deemed [ready](#) by the kernel it is handed off here. The scheduler is then in charge of state transitions when the kernel gives it appropriate notifications.

Implementation Requirements:

- Store the process in some location. It should be recognized as ready given its location but the datastructures and organization are implementation specific.

Parameters

<i>Process-Info*</i>	Process to add under scheduler's control
----------------------	--

Implements [cScheduler](#).

4.10.1.2 `ProcessInfo * cLottery::getNextToRun ()` [virtual]

After this function is called, it should be assumed by any scheduler implementation that the kernel will run the given process (unless otherwise notified). The currently running process should implicitly be considered for running next (again).

If there are processes left but all are blocked. This function should block until it receives a signal that a process is unblocked.

Implementation Requirements:

- Call `printUnblocked()` to update the trace log after a clock step.
- Block if there are > 0 processes but none can run.
- A process listed as `running` within the scheduler should be treated as ready when this method is called.
- A process returned as the 'next to run' should be marked as being in a `running` state before returning.
- If there are no more remaining processes this should return NULL.

Returns

`ProcessInfo*` Ready process to run next. May be the same as the currently running one.

Warning

Must be thread safe with block and unblock methods.

Implements `cScheduler`.

4.10.1.3 `void cLottery::initProcScheduleInfo (ProcessInfo *)` [virtual]

This method is called after the kernel has initialized all process data but before it is marked `ready`. This gives the scheduler an opportunity to initialize any scheduler specific data and assign it to the `ProcessInfo::scheduleData` member.

Implementation Requirements:

- No required actions.

Implements `cScheduler`.

4.10.1.4 `pidType cLottery::numProcesses ()` [virtual]

This returns how many processes, both running and blocked, are being handled by the scheduler.

Implementation Requirements:

- Return how many processes, running and blocked, are in the scheduler.

Returns

pidType

Implements [cScheduler](#).

4.10.1.5 void cLottery::removeProcess (ProcessInfo *) [virtual]

When a process terminates, either through normal means or an exception, the kernel will call this function to release a process from the scheduler's control. The scheduler should clean up any internal state for the process. Deallocation of process resources is left to the kernel.

Implementation Requirements:

- Scheduler should deallocate any resources it assigned to the process within the [ProcessInfo::scheduleData](#) member.
- The Scheduler should remove any pointers to the give process to avoid dereferencing a dead pointer.
- Implementations must not deallocate any memory except that mentioned above. This is handled by the kernel.
- Implementations should mark the process as [terminated](#).

Parameters

<i>Process-Info*</i>	Process to remove from scheduler
----------------------	----------------------------------

Implements [cScheduler](#).

4.10.1.6 void cLottery::setBlocked (ProcessInfo *) [virtual]

The kernel will call the scheduler with this function when the process has done an operation which causes it to block (I).

Implementation Requirements:

- Process must be marked [blocked](#) and scheduler state should be changed accordingly.
- After this call a process should not be considered for a scheduling decision

Warning

Must be thread safe. Signal handler/s may block during schedule decision.

Implements [cScheduler](#).

4.10.1.7 void cLottery::unlockProcess (ProcessInfo *) [virtual]

When a process has completed a blocking call the kernel will notify the scheduler that it should be unblocked. This operation should be very fast since it will likely be called from a signal handler.

Implementation Requirements:

- The process must be unblocked and marked [ready](#). It must be available for scheduling with the next call to `::getNextToRun`

Parameters

<i>Process-Info*</i>	Process to unblock
----------------------	--------------------

Warning

Must be thread safe. Signal handler/s may unblock during schedule decision.

Implements [cScheduler](#).

The documentation for this class was generated from the following files:


- [include/scheduler/lottery.h](#)
- [src/scheduler/lottery.cpp](#)

4.11 cMultiLevel Class Reference

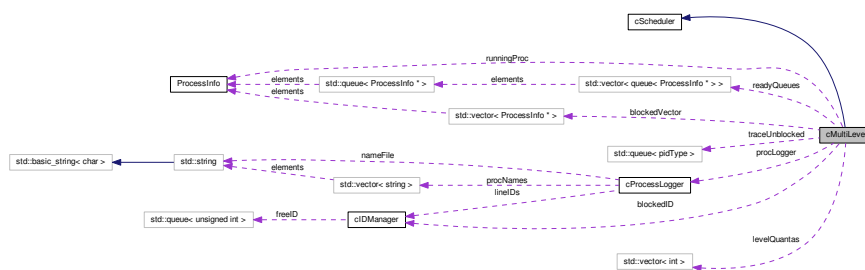
Multi-Level Queue Scheduler.

```
#include <multi_level.h>
```

Inheritance diagram for cMultiLevel:



Collaboration diagram for cMultiLevel:



Public Member Functions

- void `initProcScheduleInfo` (`ProcessInfo *`)
Initialize scheduler specific information within the `ProcessInfo` struct.
- void `addProcess` (`ProcessInfo *`)
Transfer control of process state and scheduling.
- void `setBlocked` (`ProcessInfo *`)
Set a process into a `blocked` state.
- void `unblockProcess` (`ProcessInfo *`)
Unblock a process and make it ready.
- void `removeProcess` (`ProcessInfo *`)
Remove a process from the control of the scheduler.
- `ProcessInfo *` `getNextToRun` ()
Query the scheduler for next process to run.
- pidType `numProcesses` ()

How many processes are in the scheduler.

- void [addLogger](#) (FILE *_logStream)

Add an instance of the trace logger to the scheduler.

- void [addProcLogger](#) ([cProcessLogger](#) *_procLogger)

Add instance of process logger to the scheduler.

- void [printUnblocked](#) ()

Prints at once all processes that may have become unblocked asynchronously.

- void [printLevels](#) ()

Private Attributes

- int **currentLevel**
- int **quantaUsed**
- int **totalReady**
- int **totalBlocked**
- vector< queue< [ProcessInfo](#) * > > **readyQueues**
- vector< int > **levelQuantas**
- vector< [ProcessInfo](#) * > **blockedVector**
- queue< pidType > **traceUnblocked**
- [ProcessInfo](#) * **runningProc**
- pthread_mutex_t **blockedLock**
- pthread_cond_t **allBlocked**
- [cIDManager](#) **blockedID**
- FILE * **logStream**
- [cProcessLogger](#) * **procLogger**

4.11.1 Member Function Documentation

4.11.1.1 void [cMultiLevel::addLogger](#) (FILE *) [virtual]

The kernel passes a file stream pointer for the scheduler to print trace info such as blocking/unblocking events.

Parameters

<i>FILE*</i>	File stream for trace log
--------------	---------------------------

See also

[initLog\(const char* filename\)](#)
[getStream\(\)](#)

Implements [cScheduler](#).

4.11.1.2 void cMultiLevel::addProcess (ProcessInfo *) [virtual]

After this is called, the kernel core no longer keeps track of the given process. Once the process is created and deemed [ready](#) by the kernel it is handed off here. The scheduler is then in charge of state transitions when the kernel gives it appropriate notifications.

Implementation Requirements:

- Store the process in some location. It should be recognized as ready given its location but the datastructures and organization are implementation specific.

Parameters

<i>Process-Info*</i>	Process to add under scheduler's control
----------------------	--

Implements [cScheduler](#).

4.11.1.3 void cMultiLevel::addProcLogger (cProcessLogger *) [virtual]

The kernel passes this pointer so the scheduler can update process info upon process addition, state change and termination. This updates the information presented to the top process.

Parameters

<i>cProcess-Logger*</i>	Process logger class
-------------------------	----------------------

See also

[cProcessLogger](#)

Implements [cScheduler](#).

4.11.1.4 ProcessInfo * cMultiLevel::getNextToRun () [virtual]

After this function is called, it should be assumed by any scheduler implementation that the kernel will run the given process (unless otherwise notified). The currently running process should implicitly be considered for running next (again).

If there are processes left but all are blocked. This function should block until it receives a signal that a process is unblocked.

Implementation Requirements:

- Call [printUnblocked\(\)](#) to update the trace log after a clock step.
- Block if there are > 0 processes but none can run.

- A process listed as [running](#) within the scheduler should be treated as ready when this method is called.
- A process returned as the 'next to run' should be marked as being in a [running](#) state before returning.
- If there are no more remaining processes this should return NULL.

Returns

[ProcessInfo](#)* Ready process to run next. May be the same as the currently running one.

Warning

Must be thread safe with block and unblock methods.

Implements [cScheduler](#).

4.11.1.5 void cMultiLevel::initProcScheduleInfo ([ProcessInfo](#) *) [virtual]

This method is called after the kernel has initialized all process data but before it is marked [ready](#). This gives the scheduler an opportunity to initialize any scheduler specific data and assign it to the [ProcessInfo::scheduleData](#) member.

Implementation Requirements:

- No required actions.

Implements [cScheduler](#).

4.11.1.6 pidType cMultiLevel::numProcesses () [virtual]

This returns how many processes, both running and blocked, are being handled by the scheduler.

Implementation Requirements:

- Return how many processes, running and blocked, are in the scheduler.

Returns

pidType

Implements [cScheduler](#).

4.11.1.7 void cMultiLevel::printUnblocked () [virtual]

This method is needed to avoid mixed output in the trace logger.

Implements [cScheduler](#).

Referenced by getNextToRun().

4.11.1.8 void cMultiLevel::removeProcess (ProcessInfo *) [virtual]

When a process terminates, either through normal means or an exception, the kernel will call this function to release a process from the scheduler's control. The scheduler should clean up any internal state for the process. Deallocation of process resources is left to the kernel.

Implementation Requirements:

- Scheduler should deallocate any resources it assigned to the process within the [ProcessInfo::scheduleData](#) member.
- The Scheduler should remove any pointers to the give process to avoid dereferencing a dead pointer.
- Implementations must not deallocate any memory except that mentioned above. This is handled by the kernel.
- Implementations should mark the process as [terminated](#).

Parameters

<i>Process-Info*</i>	Process to remove from scheduler
----------------------	----------------------------------

Implements [cScheduler](#).

4.11.1.9 void cMultiLevel::setBlocked (ProcessInfo *) [virtual]

The kernel will call the scheduler with this function when the process has done an operation which causes it to block (I).

Implementation Requirements:

- Process must be marked [blocked](#) and scheduler state should be changed accordingly.
- After this call a process should not be considered for a scheduling decision

Warning

Must be thread safe. Signal handler/s may block during schedule decision.

Implements [cScheduler](#).

4.11.1.10 void cMultiLevel::unblockProcess (ProcessInfo *) [virtual]

When a process has completed a blocking call the kernel will notify the scheduler that it should be unblocked. This operation should be very fast since it will likely be called from a signal handler.

Implementation Requirements:

- The process must be unblocked and marked [ready](#). It must be available for scheduling with the next call to `::getNextToRun`

Parameters

<i>Process-Info*</i>	Process to unblock
----------------------	--------------------

Warning

Must be thread safe. Signal handler/s may unblock during schedule decision.

Implements [cScheduler](#).

The documentation for this class was generated from the following files:

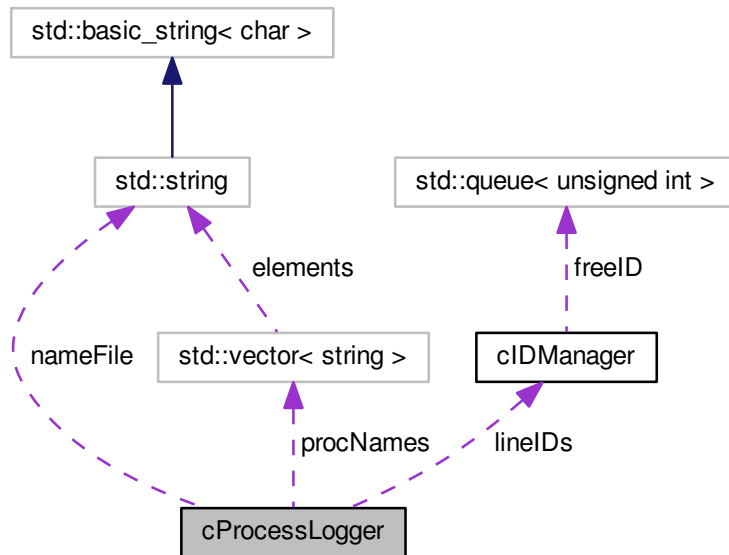
- `include/scheduler/multi_level.h`
- `src/scheduler/multi_level.cpp`

4.12 cProcessLogger Class Reference

Class specifically for logging process state information.

```
#include <process_logger.h>
```

Collaboration diagram for cProcessLogger:



Public Member Functions

- **cProcessLogger** (const char *file)
- void **addProcess** (ProcessInfo *, const char *)
- void **rmProcess** (ProcessInfo *)
- void **writeProcessInfo** (ProcessInfo *)

Private Member Functions

- void **addToVector** (FILE *)

Private Attributes

- **cIDManager** lineIDs
- string **nameFile**
- int **procLogFD**
- FILE * **procLogStream**
- int **lineSize**
- char **outputBuffer** [MAX_LINE_LENGTH]

- char **emptyBuffer** [[MAX_LINE_LENGTH](#)]
- int **previousID**
- int **listenSock**
- pthread_t **nameReqListener**
- vector< string > **procNames**
- pthread_mutex_t **logWriteLock**

Friends

- void * **nameSockFn** (void *)

4.12.1 Detailed Description

In order for monitoring programs such as top to function the kernel and its associated modules must export process information. This class logs information for each process to a file named by its pid. This is inspired by the unix /proc filesystem (although memory mapped files aren't being used). This allows the kernel to easily update only those processes which have changed.

The documentation for this class was generated from the following files:


- include/utility/[process_logger.h](#)
- src/utility/process_logger.cpp

4.13 cRoundRobin Class Reference

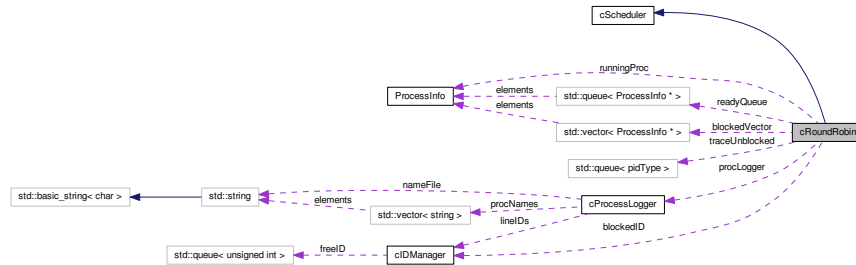
Round Robin Scheduler.

```
#include <round_robin.h>
```

Inheritance diagram for cRoundRobin:



Collaboration diagram for cRoundRobin:



Public Member Functions

- void **initProcScheduleInfo** (**ProcessInfo** *)
*Initializes a new **roundRobinInfo** struct.*
- void **addProcess** (**ProcessInfo** *)
*Adds a process to the back of the **readyQueue**.*
- void **setBlocked** (**ProcessInfo** *)
Adds a process to the blocked vector using an id created with the ID Manager.
- void **unblockProcess** (**ProcessInfo** *)
Sets the process to be ready again and adds it to the back of the queue.
- void **removeProcess** (**ProcessInfo** *)
*Sets the process state to be terminated and resets **clockTicksUsed** to 0.*
- **ProcessInfo** * **getNextToRun** ()
*If a proces is currently executing and **clockTicksUsed** is greater than the **QUANTUM**, stop the current process from executing futher and add it to the back of the ready queue.*
- pidType **numProcesses** ()
*Count up the number of processes in the **readyQueue**, processes that are blocked, and add 1 if a process is currently running.*
- void **addLogger** (FILE *_logStream)
*Assign the file pointer to the **logStream** data field.*
- void **addProcLogger** (**cProcessLogger** *_procLogger)
*Assign the **cProcessLogger** pointer to the **procLogger** data field.*
- void **printUnblocked** ()
*Removes each PID from the **traceUnblocked** queue and prints it to the trace logger.*

Private Attributes

- queue< **ProcessInfo** * > **readyQueue**
- vector< **ProcessInfo** * > **blockedVector**

- queue< pidType > **traceUnblocked**
- int **totalBlocked**
- int **clockTicksUsed**
- [ProcessInfo](#) * **runningProc**
- pthread_mutex_t **blockedLock**
- pthread_cond_t **allBlocked**
- [cIDManager](#) **blockedID**
- FILE * **logStream**
- [cProcessLogger](#) * **procLogger**

4.13.1 Member Function Documentation

4.13.1.1 [ProcessInfo](#) * **cRoundRobin::getNextToRun ()** [virtual]

If clockTicksUsed is less than the QUANTUM, increment clockTicksUsed and return the currently executing process to the kernel to keep running. If a new process needs to be scheduled, grab the top of the ready queue and increment clockTicksUsed.

Implements [cScheduler](#).

4.13.1.2 void **cRoundRobin::removeProcess ([ProcessInfo](#) * *proc*)** [virtual]

Frees the [roundRobinInfo](#) struct

Implements [cScheduler](#).

4.13.1.3 void **cRoundRobin::setBlocked ([ProcessInfo](#) * *proc*)** [virtual]

Sets runningProc to NULL and the state of the process to blocked. Sets clockTicksUsed to 0 as a new process will be picked to run. Increment the count of blocked processes.

Implements [cScheduler](#).

4.13.1.4 void **cRoundRobin::unblockProcess ([ProcessInfo](#) * *proc*)** [virtual]

Returns the blocked index id back to the ID Manager for future reuse. Decrement the count of blocked processes. Add the blocked process id to the traceUnblocked queue.

Implements [cScheduler](#).

The documentation for this class was generated from the following files:


- include/scheduler/round_robin.h
- src/scheduler/round_robin.cpp

4.14 cScheduler Class Reference

Abstract Interface for Schedulers.

```
#include <scheduler.h>
```

Inheritance diagram for cScheduler:



Public Member Functions

- virtual void `initProcScheduleInfo (ProcessInfo *)=0`
Initialize scheduler specific information within the `ProcessInfo` struct.
- virtual void `addProcess (ProcessInfo *)=0`
Transfer control of process state and scheduling.
- virtual void `setBlocked (ProcessInfo *)=0`
Set a process into a `blocked` state.
- virtual void `unlockProcess (ProcessInfo *)=0`
Unblock a process and make it ready.
- virtual void `removeProcess (ProcessInfo *)=0`
Remove a process from the control of the scheduler.
- virtual `ProcessInfo *` `getNextToRun ()=0`
Query the scheduler for next process to run.
- virtual pidType `numProcesses ()=0`
How many processes are in the scheduler.
- virtual void `addLogger (FILE *)=0`
Add an instance of the trace logger to the scheduler.
- virtual void `addProcLogger (cProcessLogger *)=0`
Add instance of process logger to the scheduler.
- virtual void `printUnblocked ()=0`
Prints at once all processes that may have become unblocked asynchronously.

4.14.1 Member Function Documentation

4.14.1.1 void cScheduler::addLogger (FILE *) [pure virtual]

The kernel passes a file stream pointer for the scheduler to print trace info such as blocking/unblocking events.

Parameters

<i>FILE*</i>	File stream for trace log
--------------	---------------------------

See also

[initLog\(const char* filename\)](#)
[getStream\(\)](#)

Implemented in [cMultiLevel](#), [cLottery](#), [cRoundRobin](#), and [cFCFS](#).

Referenced by [cKernel::cKernel\(\)](#).

4.14.1.2 void cScheduler::addProcess (ProcessInfo *) [pure virtual]

After this is called, the kernel core no longer keeps track of the given process. Once the process is created and deemed [ready](#) by the kernel it is handed off here. The scheduler is then in charge of state transitions when the kernel gives it appropriate notifications.

Implementation Requirements:

- Store the process in some location. It should be recognized as ready given its location but the datastructures and organization are implementation specific.

Parameters

<i>Process-Info*</i>	Process to add under scheduler's control
----------------------	--

Implemented in [cMultiLevel](#), [cLottery](#), [cRoundRobin](#), and [cFCFS](#).

Referenced by [cKernel::initProcess\(\)](#).

4.14.1.3 void cScheduler::addProcLogger (cProcessLogger *) [pure virtual]

The kernel passes this pointer so the scheduler can update process info upon process addition, state change and termination. This updates the information presented to the top process.

Parameters

<i>cProcess- Logger*</i>	Process logger class
------------------------------	----------------------

See also

[cProcessLogger](#)

Implemented in [cMultiLevel](#), [cLottery](#), [cRoundRobin](#), and [cFCFS](#).

Referenced by `cKernel::cKernel()`.

4.14.1.4 `ProcessInfo * cScheduler::getNextToRun ()` [pure virtual]

After this function is called, it should be assumed by any scheduler implementation that the kernel will run the given process (unless otherwise notified). The currently running process should implicitly be considered for running next (again).

If there are processes left but all are blocked. This function should block until it receives a signal that a process is unblocked.

Implementation Requirements:

- Call [printUnblocked\(\)](#) to update the trace log after a clock step.
- Block if there are > 0 processes but none can run.
- A process listed as [running](#) within the scheduler should be treated as ready when this method is called.
- A process returned as the 'next to run' should be marked as being in a [running](#) state before returning.
- If there are no more remaining processes this should return NULL.

Returns

[ProcessInfo*](#) Ready process to run next. May be the same as the currently running one.

Warning

Must be thread safe with block and unblock methods.

Implemented in [cMultiLevel](#), [cLottery](#), [cRoundRobin](#), and [cFCFS](#).

Referenced by `cKernel::boot()`.

4.14.1.5 `void cScheduler::initProcScheduleInfo (ProcessInfo *)` [pure virtual]

This method is called after the kernel has initialized all process data but before it is marked [ready](#). This gives the scheduler an opportunity to initialize any scheduler specific data and assign it to the [ProcessInfo::scheduleData](#) member.

Implementation Requirements:

- No required actions.

Implemented in [cMultiLevel](#), [cLottery](#), [cRoundRobin](#), and [cFCFS](#).

Referenced by `cKernel::initProcess()`.

4.14.1.6 `pidType cScheduler::numProcesses ()` [pure virtual]

This returns how many processes, both running and blocked, are being handled by the scheduler.

Implementation Requirements:

- Return how many processes, running and blocked, are in the scheduler.

Returns

pidType

Implemented in [cMultiLevel](#), [cLottery](#), [cRoundRobin](#), and [cFCFS](#).

Referenced by `cKernel::boot()`, and `cKernel::cKernel()`.

4.14.1.7 `void cScheduler::printUnblocked ()` [pure virtual]

This method is needed to avoid mixed output in the trace logger.

Implemented in [cMultiLevel](#), [cLottery](#), [cRoundRobin](#), and [cFCFS](#).

4.14.1.8 `void cScheduler::removeProcess (ProcessInfo *)` [pure virtual]

When a process terminates, either through normal means or an exception, the kernel will call this function to release a process from the scheduler's control. The scheduler should clean up any internal state for the process. Deallocation of process resources is left to the kernel.

Implementation Requirements:

- Scheduler should deallocate any resources it assigned to the process within the [ProcessInfo::scheduleData](#) member.

- The Scheduler should remove any pointers to the give process to avoid dereferencing a dead pointer.
- Implementations must not deallocate any memory except that mentioned above. This is handled by the kernel.
- Implementations should mark the process as [terminated](#).

Parameters

<i>Process-Info*</i>	Process to remove from scheduler
----------------------	----------------------------------

Implemented in [cMultiLevel](#), [cLottery](#), [cRoundRobin](#), and [cFCFS](#).

Referenced by `cKernel::cleanupProcess()`.

4.14.1.9 void cScheduler::setBlocked (ProcessInfo *) [pure virtual]

The kernel will call the scheduler with this function when the process has done an operation which causes it to block (I).

Implementation Requirements:

- Process must be marked [blocked](#) and scheduler state should be changed accordingly.
- After this call a process should not be considered for a scheduling decision

Warning

Must be thread safe. Signal handler/s may block during schedule decision.

Implemented in [cMultiLevel](#), [cLottery](#), [cRoundRobin](#), and [cFCFS](#).

Referenced by `cKernel::boot()`.

4.14.1.10 void cScheduler::unblockProcess (ProcessInfo *) [pure virtual]

When a process has completed a blocking call the kernel will notify the scheduler that it should be unblocked. This operation should be very fast since it will likely be called from a signal handler.

Implementation Requirements:

- The process must be unblocked and marked [ready](#). It must be available for scheduling with the next call to `::getNextToRun`

Parameters

<i>Process-Info*</i>	Process to unblock
----------------------	--------------------

Warning

Must be thread safe. Signal handler/s may unblock during schedule decision.

Implemented in [cMultiLevel](#), [cLottery](#), [cRoundRobin](#), and [cFCFS](#).

The documentation for this class was generated from the following file:

- [include/scheduler/scheduler.h](#)

4.15 fcfsInfo Struct Reference

Struct containing process info specific for FCFS scheduling.

```
#include <fcfs.h>
```

Public Attributes

- unsigned int [blockedIndex](#)

Index position in blocked vector.

The documentation for this struct was generated from the following file:


- [include/scheduler/fcfs.h](#)

4.16 kernelError Struct Reference

Struct containing kernel crash information.

```
#include <kernel.h>
```

Collaboration diagram for kernelError:



Public Attributes

- string **message**

4.16.1 Detailed Description

When the kernel crashes, important information is placed in here and then handled by the main function in `init.cpp`.

The documentation for this struct was generated from the following file:

- `include/kernel.h`

4.17 lotteryInfo Struct Reference

Struct containing process info specific for Lottery scheduling.

```
#include <lottery.h>
```

Public Attributes

- unsigned int `readyIndex`
Index position in ready vector.

- unsigned int [blockedIndex](#)
Index position in blocked vector.

The documentation for this struct was generated from the following file:

- [include/scheduler/lottery.h](#)

4.18 ProcessInfo Struct Reference

Structure for containing process state and data.

```
#include <process.h>
```

Public Attributes

- pidType **procFileLine**
- unsigned int **parent**
- unsigned int **pid**
- unsigned int **startCPU**
- unsigned int **totalCPU**
- [eProcState](#) **state**
- uint16_t **PSW**
- int **priority**
- unsigned int **PC**
- int **VC**
- char * **processText**
- void * [scheduleData](#)
Scheduler specific data.
- unsigned long **memory**

4.18.1 Detailed Description

This struture is created in the kernel when a process is initialized. It contains all process data needed for execution and for the kernel/scheduler to make desciiions on it.

4.18.2 Member Data Documentation

4.18.2.1 void* **ProcessInfo::scheduleData**

Check specific scheduler docs for the contents of this pointer. Since the process struct remains static, this gives the ability for schedulers to store their own state without the kernel having to know ahead of time.

Referenced by `cLottery::addProcess()`, `cMultiLevel::addProcess()`, `cMultiLevel::getNextToRun()`, `cFCFS::initProcScheduleInfo()`, `cRoundRobin::initProcScheduleInfo()`,

cLottery::initProcScheduleInfo(), cMultiLevel::initProcScheduleInfo(), cFCFS::removeProcess(), cRoundRobin::removeProcess(), cLottery::removeProcess(), cMultiLevel::removeProcess(), cFCFS::setBlocked(), cRoundRobin::setBlocked(), cLottery::setBlocked(), cMultiLevel::setBlocked(), cFCFS::unlockProcess(), cRoundRobin::unlockProcess(), cLottery::unlockProcess(), and cMultiLevel::unlockProcess().

The documentation for this struct was generated from the following file:

- include/process.h

4.19 roundRobinInfo Struct Reference

Struct containing process info specific for Round-Robin scheduling.

```
#include <round_robin.h>
```

Public Attributes

- unsigned int [blockedIndex](#)
Index position in blocked vector.

The documentation for this struct was generated from the following file:

- include/scheduler/round_robin.h

4.20 sMultilInfo Struct Reference

Struct containing process info specific for Multi-Level scheduling.

```
#include <multi_level.h>
```

Public Attributes

- int [blockedIndex](#)
Index position in blocked vector.
- int [level](#)
Current queue level.

The documentation for this struct was generated from the following file:


- include/scheduler/multi_level.h

Chapter 5


File Documentation

5.1 include/cpu.h File Reference

```
#include <cstdio> #include <cstdlib> #include <assert.-  
h> #include <inttypes.h> #include "utility/logger.h" Include  
dependency graph for cpu.h:
```



This graph shows which files directly or indirectly include this file:



Classes

- class `cCPU`
A class for emulating a simple cpu.

Defines

- #define `MAX_PARAMS` 2
Max number of execution parameters for any Opcode.
- #define `MAX_PARAM_SIZE` 256
Maximum size in bytes for an execution parameter.

Typedefs

- typedef unsigned int `pidType`

Enumerations

- enum `ePSW` { `PS_EXCEPTION` = 0x1, `PS_TERMINATE` = `PS_EXCEPTION` << 1, `PS_ABNORMAL` = `PS_TERMINATE` << 1, `PS_SYSCALL` = `PS_ABNORMAL` << 1, `PS_FINISHED` = `PS_SYSCALL` << 1 }
- Enumeration of Program Status Word Flags.*

5.1.1 Detailed Description

5.1.2 Define Documentation

5.1.2.1 #define MAX_PARAM_SIZE 256

Creates exception if exceeded.

5.1.3 Enumeration Type Documentation

5.1.3.1 enum ePSW

The program status word is a bit vector and this enumeration defines the meaning of particular bits. This is used in the interpretation of execution status.

Enumerator:

PS_EXCEPTION Executing process has created an exception.

PS_TERMINATE Executing process has finished.

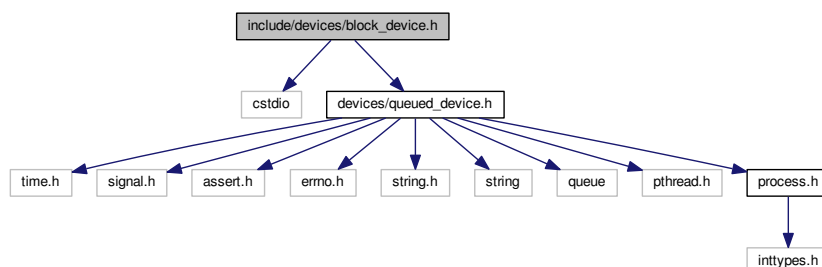
PS_ABNORMAL Process raised an exception by exiting abnormally. No 'E' instruction.

PS_SYSCALL Executing process has made a system call.


PS_FINISHED Executing process finished an instruction. No problems

5.2 include/devices/block_device.h File Reference

#include <stdio> #include "devices/queued_device.h" Include dependency graph for block_device.h:



This graph shows which files directly or indirectly include this file:



Classes

- class `cBlockDevice`

Queued block device.

Defines

- #define `BLOCKSIG` `SIGRTMIN + 1`

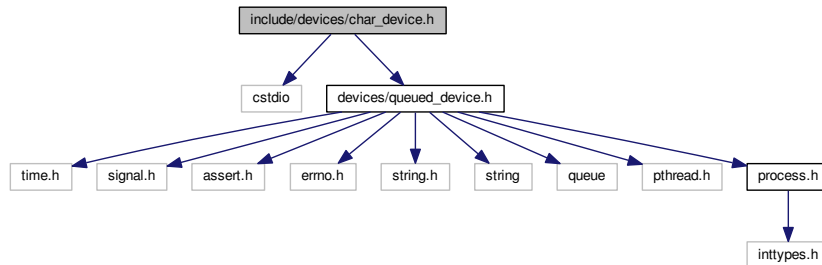
Signal generated by BlockDevice.

5.2.1 Detailed Description


5.3 include/devices/char_device.h File Reference

```
#include <stdio> #include "devices/queued_device.h" Include
```

dependency graph for char_device.h:



This graph shows which files directly or indirectly include this file:



Classes

- class `cCharDevice`
Queued char device.


Defines

- #define `CHARSIG` `SIGRTMIN + 2`
Signal generated by CharDevice.


5.3.1 Detailed Description

5.4 include/devices/clock_device.h File Reference

```
#include <errno.h> #include <cstring> #include <string> ×  
#include "devices/abstract_device.h" Include dependency graph for  
clock_device.h:
```



This graph shows which files directly or indirectly include this file:



Classes

- class [ClockDevice](#)

Device for generating repeated clock interrupts.

Defines


- #define **CLOCKID** CLOCK_REALTIME
- #define [CLOCKSIG](#) SIGRTMIN

Signal generated by [ClockDevice](#).


5.4.1 Detailed Description

5.5 include/devices/queued_device.h File Reference

```
#include <time.h> #include <signal.h> #include <assert.-  
h> #include <errno.h> #include <string.h> #include <string> ×  
#include <queue> #include <pthread.h> #include "process.-  
h" Include dependency graph for queued_device.h:
```



This graph shows which files directly or indirectly include this file:



Classes

- class [cAbsQueuedDevice](#)
Abstract class for a device which queues requests.

Defines


- #define [QD_CLOCKID](#) CLOCK_REALTIME
Type of clock used by these devices.
- #define [USEC_IN_SEC](#) 1000000
Constant used in converting times.
- #define [TIMEOUT_SCALE](#) 5
In rare occurences the signal is never received so we set a timeout.

5.5.1 Detailed Description


5.6 include/kernel.h File Reference

```
#include <vector> #include <sys/stat.h> #include <semaphore.-
h> #include "cpu.h" #include "devices/char_device.h" ×
#include "devices/block_device.h" #include "devices/clock-
```

```
_device.h" #include "scheduler/allSchedulers.h" Include dependen-
```



This graph shows which files directly or indirectly include this file:



Classes

- class `cKernel`
Core managing class for this simulated OS.
- struct `kernelError`
Struct containing kernel crash information.

Defines

- #define **DEFAULT_TIMER** 250000
Default timer for clock interrupt.
- #define **CDEVICE_SCALE** 4
Time scale for character devices relative to default clock.
- #define **BDEVICE_SCALE** 8
Time scale for block devices relative to default clock.
- #define **DEFAULT CTIMER (DEFAULT_TIMER * CDEVICE_SCALE)**

Complete time for C Devices after scaling.

- `#define DEFAULT_BTIMER (DEFAULT_TIMER * BDEVICE_SCALE)`

Complete time for B Devices after scaling.

- `#define DEFAULT_PRIORITY 5`

Default priority assigned to newly created processes.

Variables

- static const char `initProcessName []` = "main.trace"

Name of the first program to run on the system.

- static const char `traceLogFile []` = "trace.log"

Name of trace log file.

- static const char `procLogFile []` = "proc.log"

Name of log file for process info.

5.6.1 Detailed Description

5.6.2 Define Documentation

5.6.2.1 `#define DEFAULT_PRIORITY 5`

Only used if no other priority is provided.

5.6.3 Variable Documentation

5.6.3.1 static const char `initProcessName[]` = "main.trace" [static]

When the kernel object is created, this program is loaded. It is run once `cKernel::boot` is called.

Referenced by `cKernel::cKernel()`.

5.6.3.2 static const char `procLogFile[]` = "proc.log" [static]

This file is created by the process logger and it is where top gathers most of its process info.

See also

[cProcessLogger](#)

5.6.3.3 static const char traceLogFile[] = "trace.log" [static]

Per clocktick informatino is logged here.


See also

[initLog\(const char* filename\)](#)

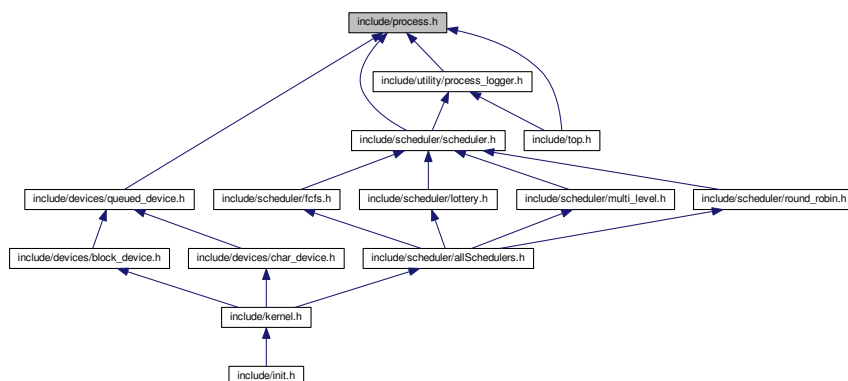
Referenced by cKernel::cKernel().

5.7 include/process.h File Reference

#include <inttypes.h> Include dependency graph for process.h:



This graph shows which files directly or indirectly include this file:



Classes

- struct [ProcessInfo](#)
Structure for containing process state and data.

Typedefs

- typedef unsigned int **pidType**

Enumerations

- enum [eProcState](#) { [ready](#), [running](#), [blocked](#), [terminated](#) }
Enumeration for process states.

5.7.1 Detailed Description

5.7.2 Enumeration Type Documentation

5.7.2.1 enum [eProcState](#)

Each values defines a current state and possible transitions.

Enumerator:

ready Process is ready to be run. Invariant State:

- Kernel has initialized it at some point
- Process should be prepared to run

Potential Transitions:

- [running](#) - Scheduler picks it to run next

running Process is currently running. A running process should implicitly be considered ready. The kernel may not notify the scheduler to transition the process to ready before asking for a scheduling decision. It is acceptable for the scheduler to make a process ready without the kernel's consent when it is being asked for a scheduling decision, assuming it was previously running.

Invariant State:

- Process is on the cpu

Potential Transitions:

- [ready](#) - Scheduler picks someone else to run
- [blocked](#) - Makes blocking system call
- [terminated](#) - Causes exception in cpu or finished normally

blocked Process is blocked and cannot run. Invariant State:

- Process is blocked (for now it can only block on I/O)

Potential Transitions:

- [ready](#) - Kernel notifies scheduler that I/O has finished

terminated Process has been terminated. It will be cleaned up soon.

Invariant State:


- Process either caused cpu exception or finished
- Process can no longer run

Potential Transitions:

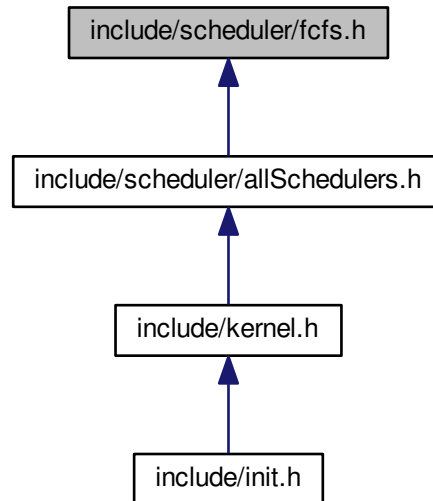
- None

5.8 include/scheduler/fcfs.h File Reference

#include <assert.h> #include "scheduler/scheduler.h" Include dependency graph for fcfs.h:



This graph shows which files directly or indirectly include this file:



Classes

- class [cFCFS](#)
First-Come-First-Serve Scheduler.
- struct [fcfsInfo](#)
Struct containing process info specific for FCFS scheduling.

Defines


- #define **DEF_BLOCK_VEC_SIZE** 4

5.8.1 Detailed Description

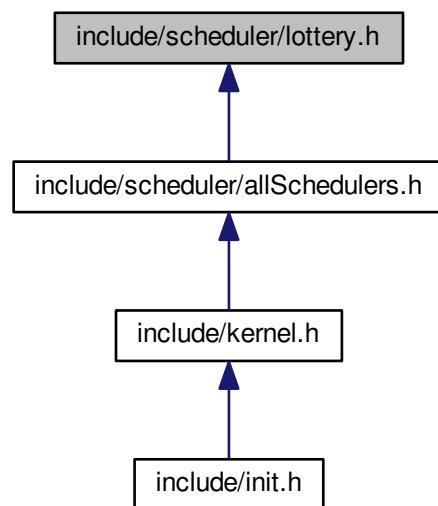
5.9 include/scheduler/lottery.h File Reference

```
#include <assert.h> #include <pthread.h> #include <time.-
h> #include "scheduler/scheduler.h" #include "utility/id.-
```


h " Include dependency graph for lottery.h:



This graph shows which files directly or indirectly include this file:



Classes

- class [cLottery](#)

Lottery Scheduler.

- struct [lotteryInfo](#)

Struct containing process info specific for Lottery scheduling.


Defines

- #define **DEF_BLOCK_VEC_SIZE** 4

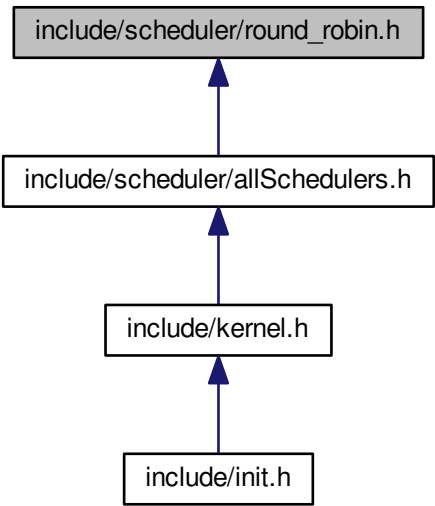
5.9.1 Detailed Description

5.10 include/scheduler/round_robin.h File Reference

#include <assert.h> #include "scheduler/scheduler.h" Include dependency graph for round_robin.h:



This graph shows which files directly or indirectly include this file:



Classes

- class [cRoundRobin](#)

Round Robin Scheduler.

- struct [roundRobinInfo](#)

Struct containing process info specific for Round-Robin scheduling.


Defines

- #define **DEF_BLOCK_VEC_SIZE** 4
- #define **QUANTUM** 4

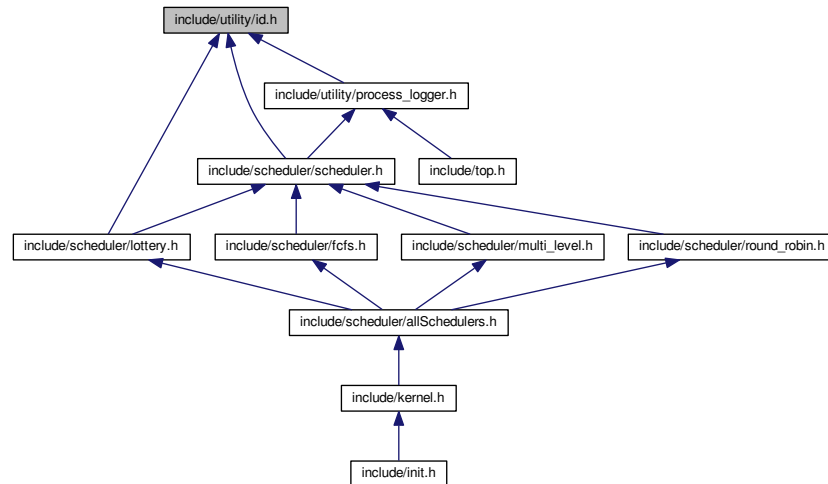
5.10.1 Detailed Description

5.11 include/utility/id.h File Reference

```
#include <stdio> #include <climits> #include <string> ×  
#include <queue> Include dependency graph for id.h:
```



This graph shows which files directly or indirectly include this file:



Classes

- class [cIDManager](#)


A class for managing unique IDs.

5.11.1 Detailed Description

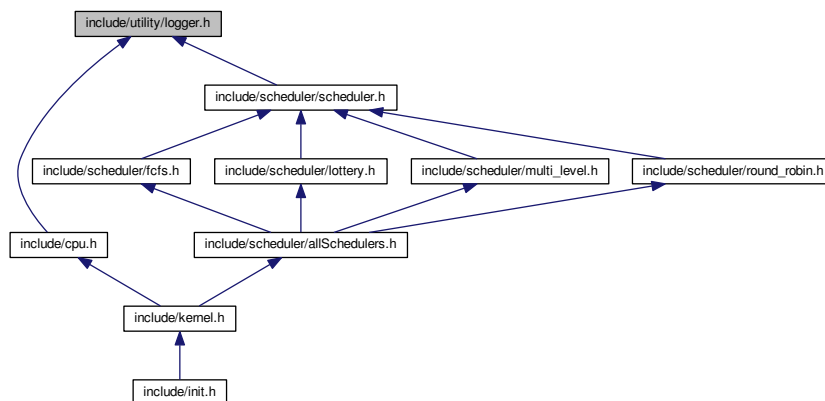
5.12 include/utility/logger.h File Reference

```
#include <assert.h> #include <iostream> #include <fstream> ×
```

Include dependency graph for logger.h:



This graph shows which files directly or indirectly include this file:




Functions

- FILE * [initLog](#) (const char *filename)
Initialize a trace log at filename.
- void [closeLog](#) ()
Close the file stream for the trace log.
- FILE * [getStream](#) ()
Get the file stream to write to.

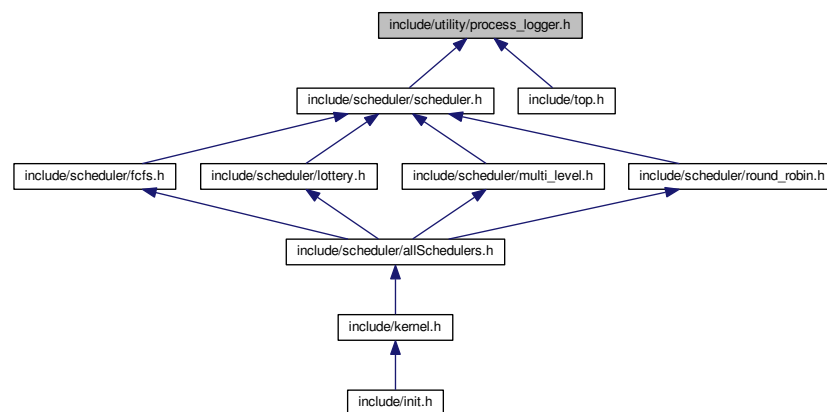
5.12.1 Detailed Description

5.13 include/utility/process_logger.h File Reference

```
#include <stdio.h> #include <stdlib.h> #include <unistd.-
h> #include <fcntl.h> #include <assert.h> #include <errno.-
h> #include <string.h> #include <pthread.h> #include
<sys/types.h> #include <sys/socket.h> #include <sys/un.-
h> #include "process.h" #include "utility/id.h" Include depen-
dency graph for process_logger.h:
```



This graph shows which files directly or indirectly include this file:



Classes

- class [cProcessLogger](#)
Class specifically for logging process state information.

Defines

- #define [MAX_LINE_LENGTH](#) 45

Max line length for a process entry in the log file.

Enumerations

- enum **pivotType** { **pivotMiddle**, **pivotRandom** }

Functions

- template<typename T >
void **QuickPartition** (std::vector< T > items, int(*sort_fn)(T, T), int left, int right, int pivot=pivotMiddle)
- template<typename T >
void **QuicksortVector** (std::vector< T > items, int(*sort_fn)(T, T), pivotType p-Type, int start, int end)

Variables

- static const char **procNameReq** [] = "proc.log.req"
Name of unix socket file for processes like top to request process names.
- static const char **outputFormat** [] = "%u %d %d %d %d"
Format for process info in the log file.
- static const char **requestError** [] = "INVALID_ID"
Return value when a invalid process ID was requested.

5.13.1 Detailed Description

5.13.2 Variable Documentation

5.13.2.1 static const char **outputFormat**[] = "%u %d %d %d %d" [static]

Scanf is used to print the information in this format to a buffer. Then this buffer is padded to fill **MAX_LINE_LENGTH** and then it is written out to the appropriate line in the file.

5.13.2.2 static const char **procNameReq**[] = "proc.log.req" [static]

Since process trace names are variable, their names are stored in a string objects and then served upon request to this socket. While filenames do have a max size, it is more efficient to do it this way when you consider that most trace file names will not be near the max.

5.13.2.3 static const char **requestError**[] = "INVALID_ID" [static]

If another process requests an invalid ID on the request socket, this is the corresponding message.