

Project 3: File Systems

Due: 10 PM April 17, 2012

Introduction

In this project, you will add features to an existing file system simulator. The file system simulator shows the inner workings of a UNIX V7 file system. The simulator reads or creates a file which represents the disk image, and keeps track of allocated and free blocks using a bit map.

The File System Simulator is a collection of C++ classes which simulate the file system calls available in a typical Unix-like operating system. The "Kernel" class contains methods (functions) like "creat()", "open()", "read()", "write()", "close()", etc., which read and write blocks in an underlying file in much the same way that a real file system would read and write blocks on an underlying disk device.

In addition to the "Kernel" class, there are a number of underlying classes to support the implementation of the kernel. The classes FileSystem, IndexNode, DirectoryEntry, SuperBlock, Block, BitBlock, FileDescriptor, and Stat contain all data structures and algorithms which implement the simulated file system.

Also included are a number of sample programs which can be used to operate on a simulated file system. The programs "ls", "cat", "mkdir", "mkfs", etc., perform file system operations to list directories, display files, create directories, and create (initialize) file systems. These programs illustrate the various file system calls and allow the user to carry out various read and write operations on the simulated file system.

As mentioned above, there is a backing file for our simulated file system. A "dump" program is included with the distribution so that you can examine this file, byte-by-byte. Any dump program may be used (e.g., the "od" program in Unix); we include this one which is simple to use and understand, and can be used with any operating system.

There are a number of ways you can use the simulator to get a better understanding of file systems. You can

- use the provided utility programs (`mkfs`, `mkdir`, `ls`, `cat`, etc.) to perform operations on the simulated file system and use the `dump` program to examine the underlying file and observe any changes,
- examine the sample utility programs to see how they use the system call interface to perform file operations,
- enhance the sample utility programs to provide additional functionality,
- write your own utility programs to extend the functionality of the simulated file system, and
- modify the underlying Kernel and other implementation classes to enhance the simulator

In the sections which follow, you will learn what you need to know to perform each of these activities. And in this project, you will do several of these.

Using File System Simulator Programs

Using `mkfs`

The `mkfs` program creates a file system backing file. It does this by creating a file whose size is specified by the block size and number of blocks given. It writes the superblock, the free list blocks, the inode blocks, and the data blocks for a new file system. Note that it will overwrite any existing file of the name specified, so be careful when you use this program.

This program is similar to the "mkfs" program found in Unix-like operating systems.

The general format for the `mkfs` command is:

```
./mkfs file-name block-size blocks
```

where

file-name is the name of the backing file to create (e.g., `filesys.dat`). Note that this is the name of a real file, not a file in simulator. This is the file that the simulator uses to simulate the disk device for the simulated file system. This may be any valid file name in your operating system environment.

block-size is the block size to be used for the file system (e.g., 256). This should be a multiple of the index node (i-node) size (usually 64) and the directory entry size (usually 16). Modern operating systems usually use a size of 1024, or 512 bytes. We use 128 or 256 byte block sizes in many of our examples so that you can quickly see what happens when directories grow beyond one block. This should be a decimal number not less than 64, but less than 32768.

blocks is the number of blocks to create in the file system (e.g., 40). This number includes any blocks that may be used for the superblock, free list management, inodes, and data blocks. We use a relatively small number here so that you can quickly see what happens if you run out of disk space. This can be any decimal number greater than 3, but not greater than $2^{24} - 1$ (the maximum number of blocks), although you may not have sufficient space to create a very large file.

For example, the command:

```
./mkfs filesys.dat 256 40
```

will create (or overwrite) a file "filesys.dat" so that it contains 40 256-byte blocks for a total of 10240 bytes.

The output from the command should look something like this:

```
block_size: 256
```

```
blocks: 40
```

```
super_blocks: 1
```

```
free_list_blocks: 1
```

```
inode_blocks: 8
```

```
data_blocks: 30
```

```
block_total: 40
```

From the output you can see that one block is needed for the superblock, one for free list management, eight for index nodes, and the remaining 30 are available for data blocks.

Why is there 1 block for free list management? Note that 30 blocks require 30 bits in the free list bitmap. Since $256 \text{ bytes/block} * 8 \text{ bits/byte} = 2048 \text{ bits/block}$, clearly one bitmap block is sufficient to track block allocation for this file system.

Why are there 8 blocks for index nodes? Note that 30 blocks could result in 30 inodes if many one-block files or directories are created. Since each inode requires 64 bytes, only 4 will fit in a block. Therefore, 8 blocks are set aside for up to 32 inodes.

NOTE: The other utilities assume that `filesys.dat` has been initialized using `mkfs` and is located in the same directory as the utility's binary. Using the other utilities without a properly initialized `filesys.dat` will result in undefined behavior.

Using `mkdir`

The `mkdir` program can be used to create new directories in our simulated file system. It does this by creating the file specified as a directory file, and then writing the directory entries for "." and ".." to the newly created file. Note that all directories leading to the new directory must already exist.

This program is similar to the "mkdir" command in Unix-like and MS-DOS-related operating systems.

The general format for the `mkdir` command is:

```
./mkdir directory-path
```

where

directory-path is the path of the directory to be created (e.g., "/root", or "temp", or "../home/jon/os/filesys"). If *directory-path* does not begin with a "/", then it is appended to the path name for working directory for the default process.

For example, the command:

```
./mkdir /home
```

creates a directory called "home" as a subdirectory of the root directory of the file system.

Similarly, the command:

```
./mkdir /home/jon
```

creates a directory called "jon" as a subdirectory of the "home" directory, which is presumed to already exist as a subdirectory of the root directory of the file system.

Using `ls`

The `ls` program is used to list information about files and directories in our simulated file system. For each file or directory name given it displays information about the files named, or in the case of directories, for each file in the directories named.

This program is similar to the "ls" command in Unix-like operating systems.

The general format for the `ls` command is:

```
./ls path-name ...
```

where

path-name is a space-separated list of one or more file or directory path names.

For example, the command:

```
./ls /home
```

lists the contents of the "/home" directory. For each file in the directory, a line is printed showing the name of the file or subdirectory, and other pertinent information such as size.

The output from the command should look something like this:

```
/home:
```

```
1          48 .
0          48 ..
2          32 jon
```

```
total files: 3
```

In this case we see that the "/home" directory contains entries for ".", "..", and "jon".

Using `tee`

The `tee` program reads from standard input and writes whatever is read to both standard output and the named file. You can use this program to create files in our simulated file system with content created in the operating system environment.

This program is similar to the "tee" command found in many Unix-like operating systems.

The general format for the `tee` command is:

```
./tee file-path
```

where

file-path is the name of a file to be created in the simulated file system. If the named file already exists, it will be overwritten.

For example:

```
echo "howdy, podner" | ./tee /home/jon/hello.txt
```

causes the single line "howdy, podner" to be written to the file "/home/jon/hello.txt".

The output from the command is:

```
howdy, podner
```

which you should note was the same as the input sent to the `tee` program by the "echo" command.

Note that the "|" (pipe) is almost always used with the `tee` program. Users of Unix-like operating systems will find the "echo", and "cat" commands useful to produce input for the pipe to `tee`..

If you wish to simply enter text directly to a file, then you may use `tee` directly (i.e., without the pipe). Users of Unix-like operating systems will need to use CTRL-D to signal the end of input.

Using `cp`

The `cp` program allows you to copy the contents from one file to another in our simulated file system. If the destination file already exists, it will be overwritten.

The general format of the "cp" command is:

```
./cp input-file-name output-file-name
```

where

input-file-name is the path-name for the file to be copied (i.e., the *source* file, and

output-file-name is the path-name for the file to be created (i.e., the *target* file.

For example:

```
./cp /home/jon/hello.txt /home/jon/greeting.txt
```

creates a new file "/home/jon/greeting.txt" by copying to it the contents of file "/home/jon/hello.txt".

Using `cat`

The `cat` program reads the contents of a named file and writes it to standard output. The `cat` program is generally used to display the contents of a file.

The general format of the `cat` command line is:

```
./cat file-name
```

where

file-name is the name of the file from which data are to be read for writing to standard output.

For example:

```
./cat /home/jon/greeting.txt
```

causes the file `/home/jon/greeting.txt` to be read, the contents of which are written to standard output.

In this case, the output from the program might look something like this:

```
howdy, podner
```

Dumping the File System

While you are working with the file system simulator, you may wish to dump the contents of the backing file to see if it contains what you *think* it contains. The `dump` program shows the contents of a file in the operating environment, one byte at a time, in various formats (hexadecimal, decimal, ASCII).

Note that `dump` dumps the contents of a real file, not a file in our simulated file system.

The general format of the `dump` command line is

```
./dump file-name
```

where

file-name

is the name of the file to be dumped. This should generally be the name of the backing file for the file system simulator (e.g., `filesys.dat`).

The general format of the dump output is

```
addr hex dec asc
```

where

addr is the decimal address of the byte,

hex is the hexadecimal value of the byte,

dec is the decimal value of the byte, and

asc is the corresponding ASCII character if the value is between 33 and 127 (decimal).

Each line of `dump` output corresponds to a single byte in the file. To keep the listing brief, `dump` only displays non-zero bytes from the input file.

For example

```
./dump filesys.dat | more
```

causes the contents of the file `filesys.dat` to be displayed, one line per byte. The `"| more"` causes you to be prompted for each page of the output.

The first page of the output should look something like this:

```
0 1 1
5 28 40 (
9 1 1
13 2 2
17 a 10
256 1f 31
512 40 64 @
515 3 3
523 30 48 0
527 ff 255
528 ff 255
529 ff 255
530 ff 255
531 ff 255
532 ff 255
533 ff 255
534 ff 255
535 ff 255
536 ff 255
537 ff 255
538 ff 255
539 ff 255
540 ff 255
541 ff 255
```

You should notice, for example, that the first block (the super block) contains a few numeric values corresponding to the block size (the 1 in the 0 byte means 256), number of blocks, etc. The second block (starting at byte 256) contains a few bits that are set, indicating that the first few blocks are allocated. The third block (starting at 512) contains a few index nodes; the FF/255 values indicate that a direct block is unallocated. A little further down you will see ".", and ".." for the directory entries for the root file system, and other data blocks.

Writing File System Simulator Programs

Writing programs that use the File System Simulator requires the use of the `Kernel` class, and may involve the use of the classes `Stat` and `DirectoryEntry`. Each file system simulator program must call `Kernel.initialize()` before calling any of the other `Kernel` methods. If you're writing ordinary programs that use the standard file system calls, you should not need to reference any other classes.

These three classes are described briefly here.

Kernel

sets up the simulator environment and defines all the system calls. This class defines: the method `initialize()`, which is used to initialize the file system simulator; the `creat()`, `open()`, `read()`, `write()`, `close()`, and other methods which simulate the work of a file system; and constants like `EBADF`, `S_IFDIR`, and `O_RDONLY` which are used to represent parameter or return values for the system calls. All the methods and fields of `Kernel` are static; you do not instantiate a `Kernel` object. For examples, see any of the sample programs (i.e., `cat.cc`, `cp.cc`, `ls.cc`, etc.)

Stat

is a data structure that represents information about a file or directory. This intends to faithfully represent the Unix `stat` struct. You may reference fields within a `stat` object directly (e.g., `stat.st_ino`), or using accessor/mutator methods (e.g., `stat.getIno()` or `stat.setIno()`). `Stat` objects are updated by the methods `Kernel.stat()` and `Kernel.fstat()`. For examples, see `mkdir.cc`.

DirectoryEntry

is a data structure that represents a single record in a directory file. This intends to faithfully represent a Unix `dirent` struct. It contains an index node number and a file name. You may reference the fields directly (e.g., `dirent.d_ino`), or using accessor/mutator methods (e.g., `dirent.getIno()` or `dirent.setIno()`). `DirectoryEntry` objects are updated by the method `Kernel.readdir()`. For examples, see `mkdir.cc` and `ls.cc`.

For more information about Unix system calls and the `stat` and `dirent` structs, refer to a Unix system manual. Users of Unix-like systems may find the commands "man -S 2 creat", "man -S 2 open", etc. to be helpful.

All programs that use the File System Simulator should adhere to the following guidelines:

- Invoke the method `Kernel.initialize()` before any other File System Simulator calls.
- Use `Kernel.exit()` when you wish to terminate processing in your program.
- Check for errors after each system call (e.g., `creat()`, `open()`, `read()`, `write()`, etc.). Nearly all the system calls return -1 if an error occurs.
- Use `Kernel.perror()` to print the message associated with an error.
- Use `Kernel.getErrno()` to determine which error occurred, if needed. Note that in standard Unix programs you would reference the static process variable "errno".

For examples, take a look at the following sample programs in the distribution:

- `cat.cc`
- `cp.cc`
- `ls.cc`
- `mkdir.cc`
- `tee.cc`

Collectively, these sample programs invoke all of the core methods (system calls) of the file system simulator.

The File System Classes

The following are the *internal* classes for the file system simulator:

`BitBlock`

is a data structure that views a device block as a sequence of bits. The methods `setBit()`, `resetBit()`, and `isBitSet()` are used to set, reset, or check a bit in the block. This structure is used to implement bitmaps, and is used by the file system simulator to track allocated and free data blocks in the file system. `BitBlock` extends `Block`.

`Block`

is a data structure that views a device block as a sequence of bytes. The field `bytes` is an array of `byte`, and is directly accessible. Included are methods to `read()` and `write()` the block to `aRandomAccessFile`, which simulate the action of reading or writing a device block.

`FileDescriptor`

is a structure and collection of methods that represent an open file. It includes a number of `get` and `set` methods for various tidbits of information about the open file, and provides `readBlock` and `writeBlock()` methods for reading and writing the blocks of the file.

`FileSystem`

is a structure and collection of methods that represent an open (mounted) file system. It includes a few `get` and `set` methods for various fields about the file system, but more importantly, includes methods to `open()` the file behind the file system, to `read()` and `write()` blocks of the device, to manage blocks (`allocateBlock()` and `freeBlock()`) and to manage inodes (`allocateIndexNode()`). In general, Kernel methods should call `FileSystem` methods when they want to read or write data in the file system.

`IndexNode`

is a structure and collection of methods for representing an index node. This is meant to reflect the exact structure on disk for an index node. It includes `get` and `set` methods for each of the fields in the index node. Also included are `read()` and `write()` methods which are used to copy data to and from byte arrays (not disk files).

`ProcessContext`

is a structure and collection of methods to represent a process. This is where the simulator stores the uid, gid, umask, dir, and other information for the current process. It includes `get` and `set` methods for each of the fields in a process.

SuperBlock

is a structure and collection of methods for representing the superblock on the disk. In our implementation, the superblock contains information about the block size, number of blocks, offsets to the first block of the free list, inode block, and data block areas of the device. It includes `get` and `set` methods for each of the fields in the superblock. Also included are methods to `read()` and `write()` the superblock.

Objectives

Study the code and utilities to gain a basic understanding of the file system code, before adding the features described below.

1. Enhance the file system simulator to include a new method, `Kernel.link()`, which, given two path names, creates the second path as a (hard) link to the first path. `link()` should find the inode number for the first file, and then write a directory entry for the second path which references the same index node. Don't forget to increment `nlink` on the index node. To test your new method, write a new program, `ln.cc`, which, given two path names, performs the `link()` operation. Assume that creating a link to a directory is not allowed.
2. Enhance the file system simulator to include a new method, `Kernel.unlink()`, which, given the name of a file, removes the directory entry for that file and decrements `nlink` for the index node. If `nlink` is decremented to zero, free all the blocks of the file. If the file is currently open by any process, mark the file so that the blocks will be freed when the file is closed by the last process. To test your new method, write a new program `rm.cc` which accept the names of files to be unlinked. Assume that unlinking a directory is not allowed.
3. Enhance the file system simulator to support indirect blocks, and double- and triple-indirect blocks.
4. Write a program called `fsck.cc` which checks a file system for internal consistency. It should verify that all the inodes mentioned in directory entries have the correct number of `nlinks`, and that all blocks mentioned in the inodes are marked as allocated blocks, and all blocks NOT mentioned in the inodes are marked as free blocks. At a minimum you should list any problems to standard out. For simplicity, you can check only direct blocks.
Repair the errors during an `fsck`. Decide on a reasonable fix, and then ask the end-user to authorize your repair. Also – check all blocks (at all levels of indirection).

Extra Credit

1. Repair the errors during an fsck. Decide on a reasonable fix, and then ask the end-user to authorize your repair. Also – check all blocks (at all levels of indirection).
2. Write a program `find.cc` which, given a path name, checks to see if the path exists, and if so lists that path name and all files in all directories (and sub-directories, and sub-sub-directories, etc.) under it, one path name per line. For example, `./find /home` might produce the following output:

`/home`

`/home/foo`

`/home/foo/bar.txt`

`/home/foo/foo.txt`

`/home/jon`

`/home/jon/homer`

`/home/jon/homer/odyssey.txt`

`/home/jon/homer/iliad.txt`

`/home/jon/virgil`

`/home/jon/virgil/aeneid.txt`

`/home/jon/virgil/eclogues.txt`

`/home/jon/virgil/georgics.txt`

under the right circumstances, of course. Hint: Your program may include a recursive method or an array for keeping track of each directory as you open it. What is the maximum directory tree depth to which your program will work?

3. Implement symbolic links.

You must completely implement a feature to gain extra credit.

Submission

- Zip up/Tar your code and submit it online.
- Projects **MUST** include a `readme.txt` file in your assignment which has the following information:
 - Your names;
 - Your student IDs;
 - How to build your program (ideally via “make” or “make all”)

- How to run your program;
- No need to submit a paper copy. Save trees instead.